

# **CSCI 3753**

# **Operating Systems**

## **Device Management**

**Lecture Notes By**

**Shivakant Mishra**

**Computer Science, CU-Boulder**

**Last Update: 01/22/14**

# Bootstrapping the OS

- Multi-stage procedure:
  1. Power On Self Test (POST) from ROM
    - Check hardware, e.g. CPU and memory, to make sure it's OK
  2. BIOS (Basic Input/Output System) looks for a device to boot from...
    - May be prioritized to look for a USB flash drive or a CD/DVD-ROM drive before a hard disk drive
    - Can also boot from network

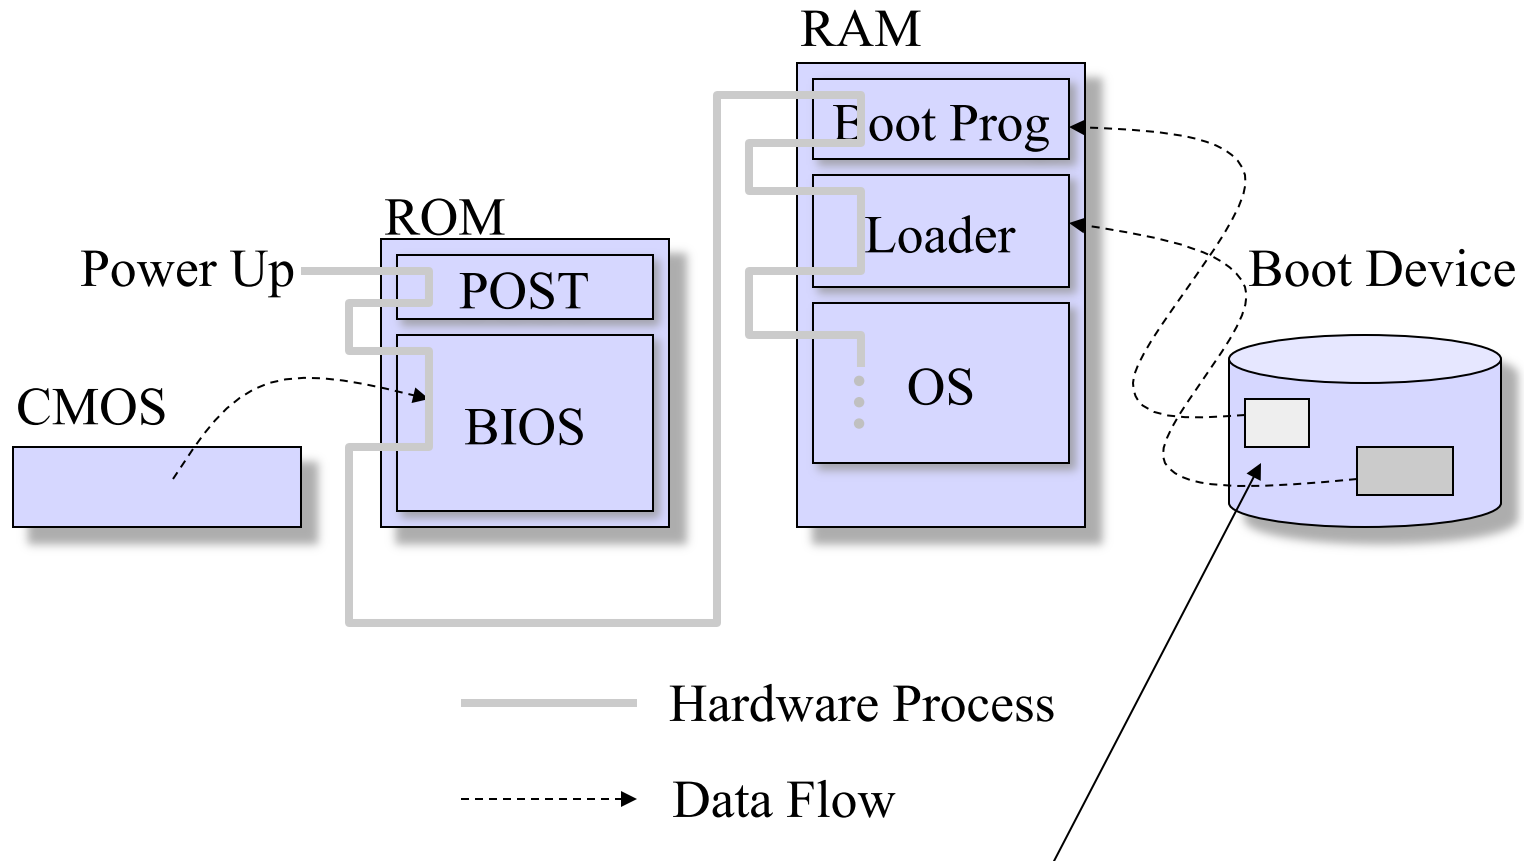
# Bootstrapping the OS (2)

- Multi-stage procedure: (continued)
  3. BIOS finds a hard disk drive to boot from
    - Looks at Master Boot Record (MBR) in sector 0 of disk
    - Only 512 bytes long (Intel systems), contains primitive code for later stage loading and a partition table listing an active partition, or the location of the bootloader

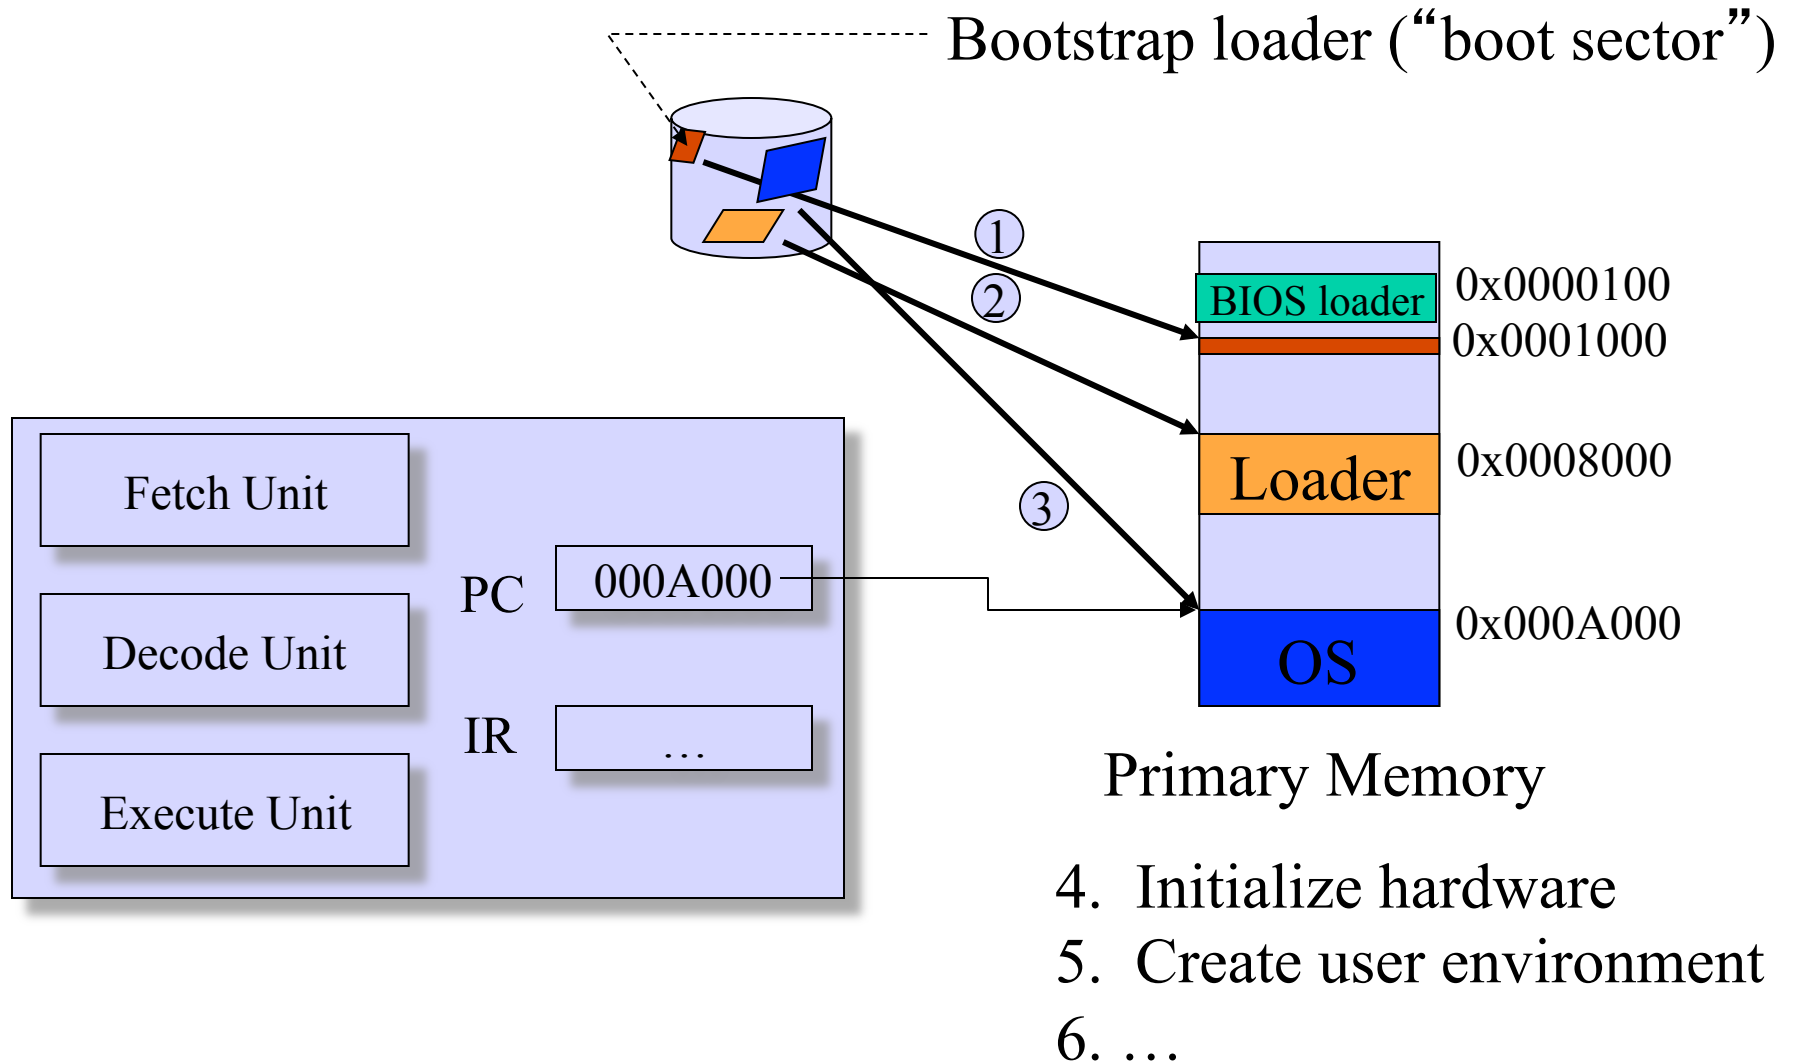
# Bootstrapping the OS (3)

- Multi-stage procedure: (continued)
  4. Primitive loader then loads the secondary stage bootloader
    - Examples of this bootloader include LILO (Linux Loader), and GRUB (Grand Unified Bootloader)
    - Can select among multiple OS' s (on different partitions) – i.e. dual booting
    - Once OS is selected, the bootloader goes to that OS' s partition, finds the boot sector, and starts loading the OS' s kernel

# Intel System Initialization



# Bootstrapping Example



# Device Manager

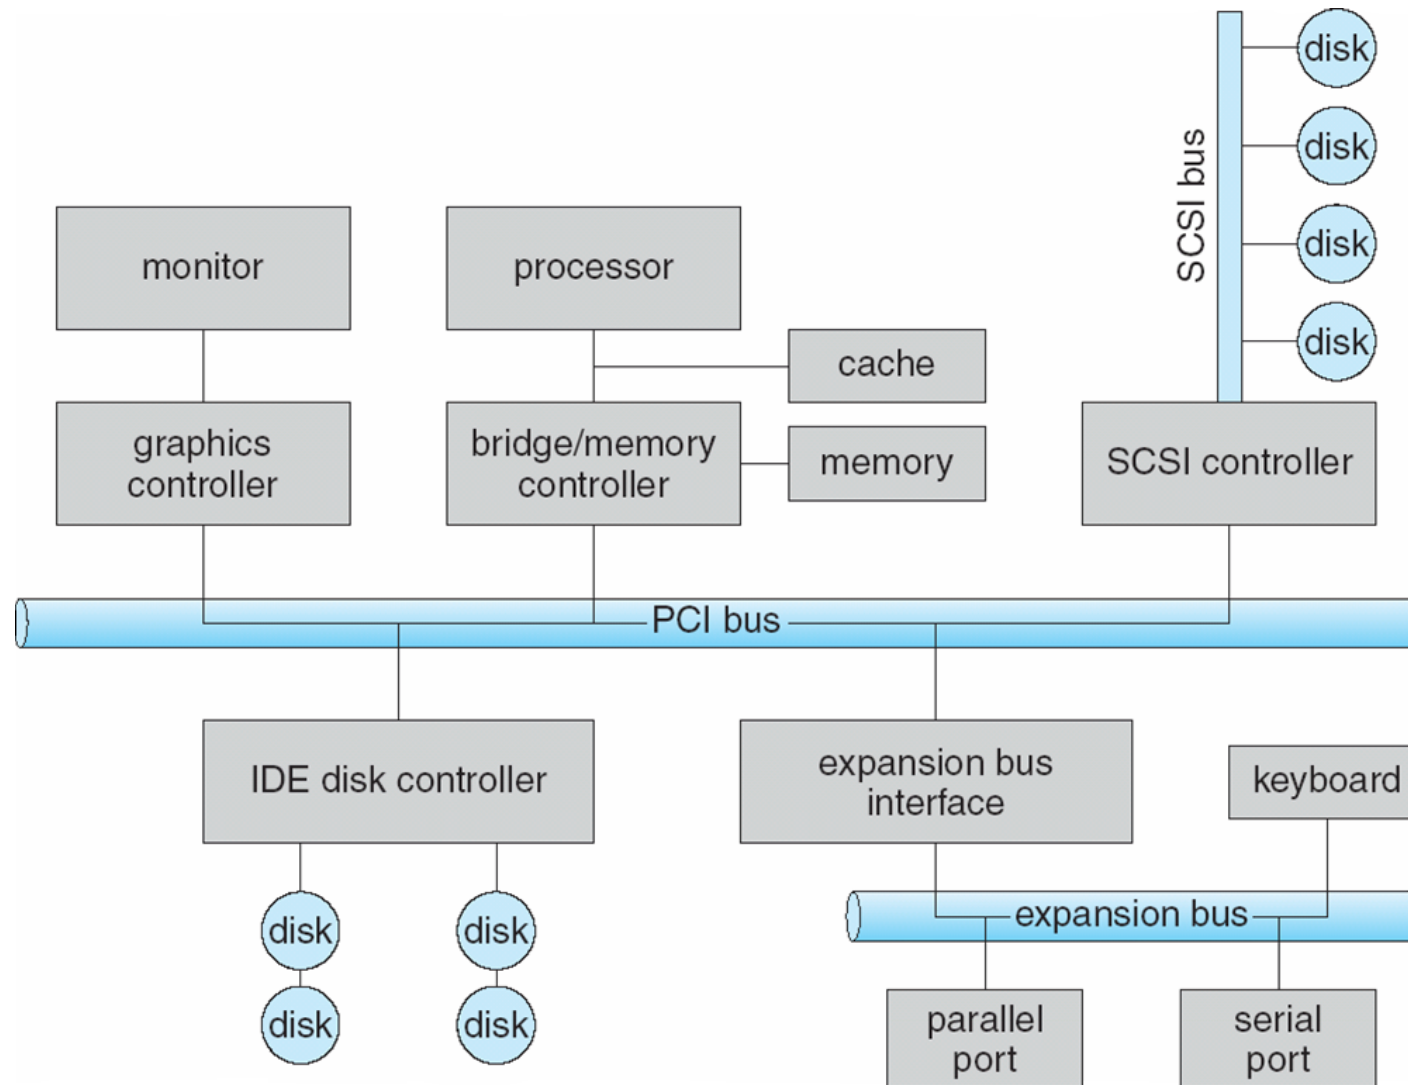
- Controls the operation of I/O devices
  - Issue I/O commands to the devices
  - Catch interrupts
  - Handle errors
  - Provide a simple and easy-to-use interface
    - Device independence: same interface for all devices.

# Device Characteristics

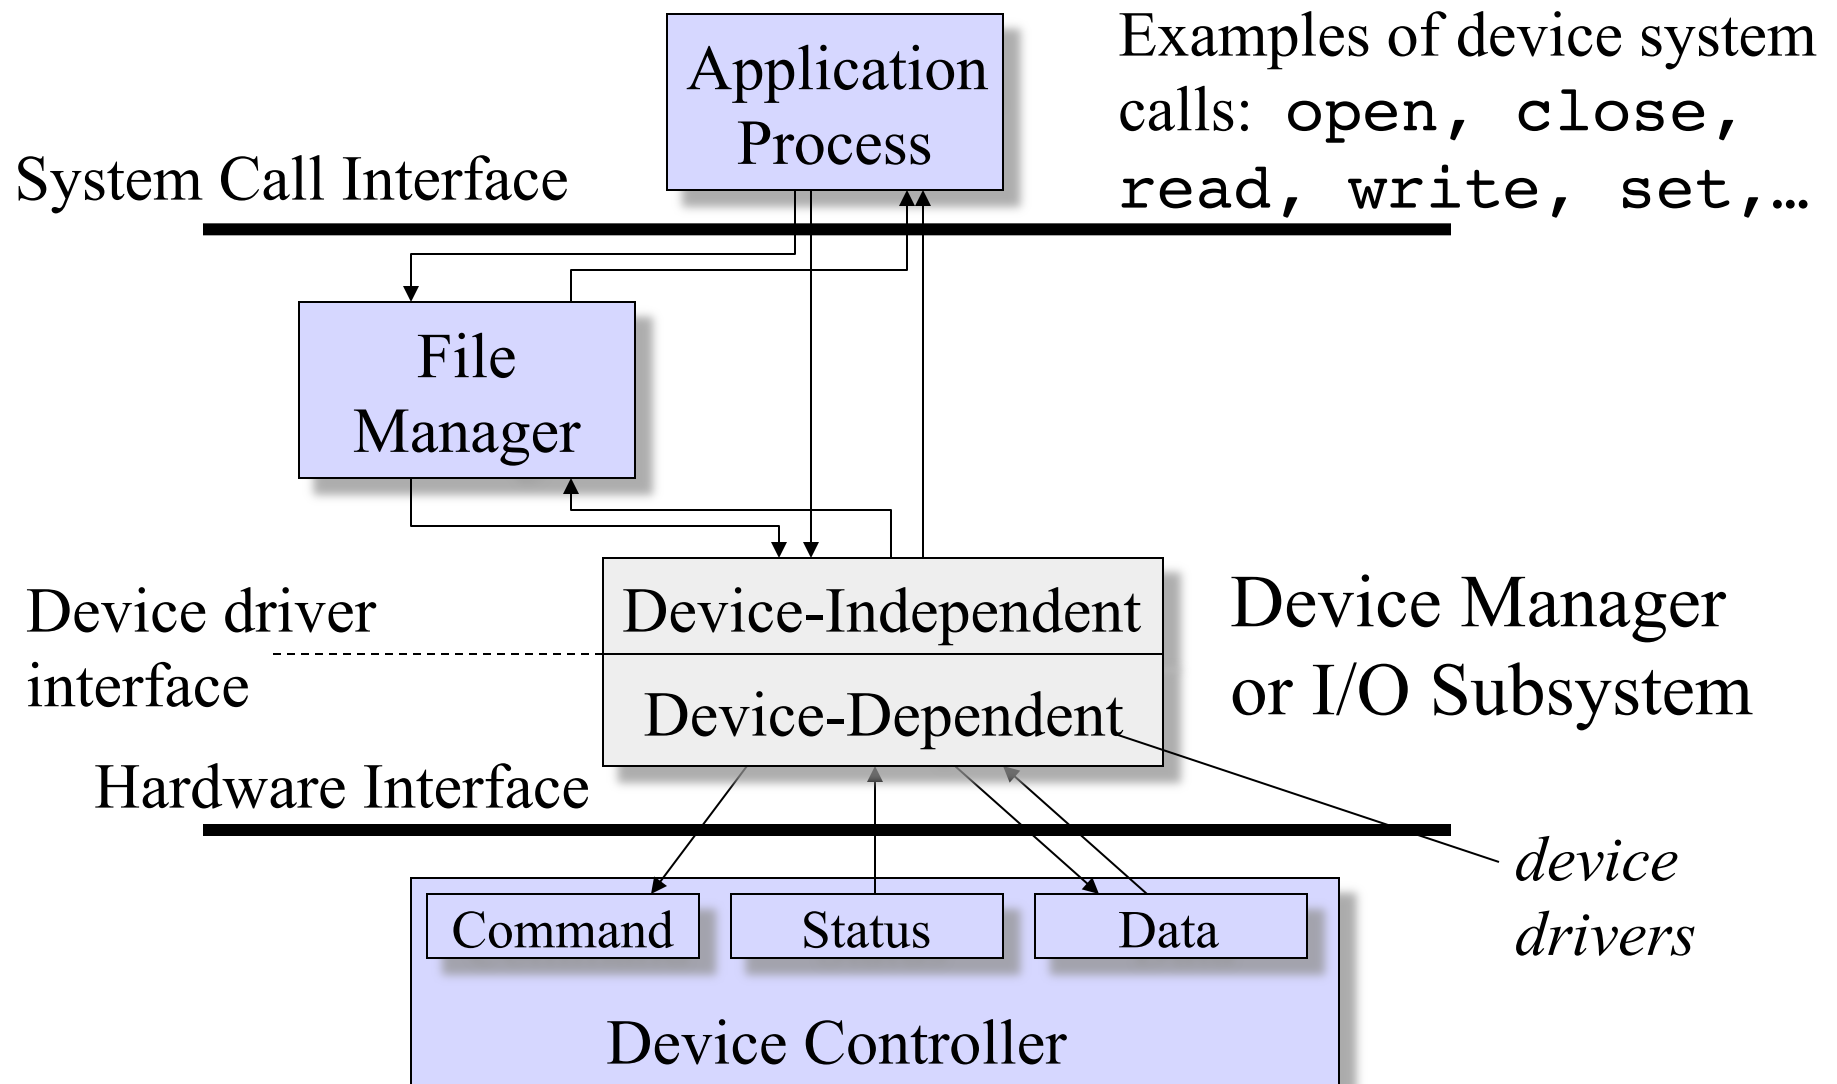
- I/O devices consist of two high-level components
  - Mechanical component
  - Electronic component: device controllers
- OS deals with device controllers



# A Typical PC Bus Structure



# Device Management Organization



Operating Systems: A Modern Perspective


# Device System Call Interface

- Create a simple standard interface to access most devices:
  - Every I/O device driver should support the following: open, close, read, write, set (ioctl in UNIX), stop, etc.
  - Block vs character
  - Sequential vs direct/random access
  - Blocking versus Non-Blocking I/O
    - blocking system call: process put on wait queue until I/O completes
    - non-blocking system call: returns immediately with partial number of bytes transferred, e.g. keyboard, mouse, network sockets
  - Synchronous versus asynchronous
    - asynchronous returns immediately, but at some later time, the full number of bytes requested is transferred

# Device Drivers

- Support the device system call interface functions `open`, `read`, `write`, etc. for that device
- Interact directly with the device controllers
  - Know the details of what commands the device can handle, how to set/get bits in device controller registers, etc.
  - Are part of the device-dependent component of the device manager
- Control flow:
  - An I/O system call traps to the kernel, invoking the trap handler for I/O (the device manager), which indexes into a table using the arguments provided to run the correct device driver

# Device Controller States

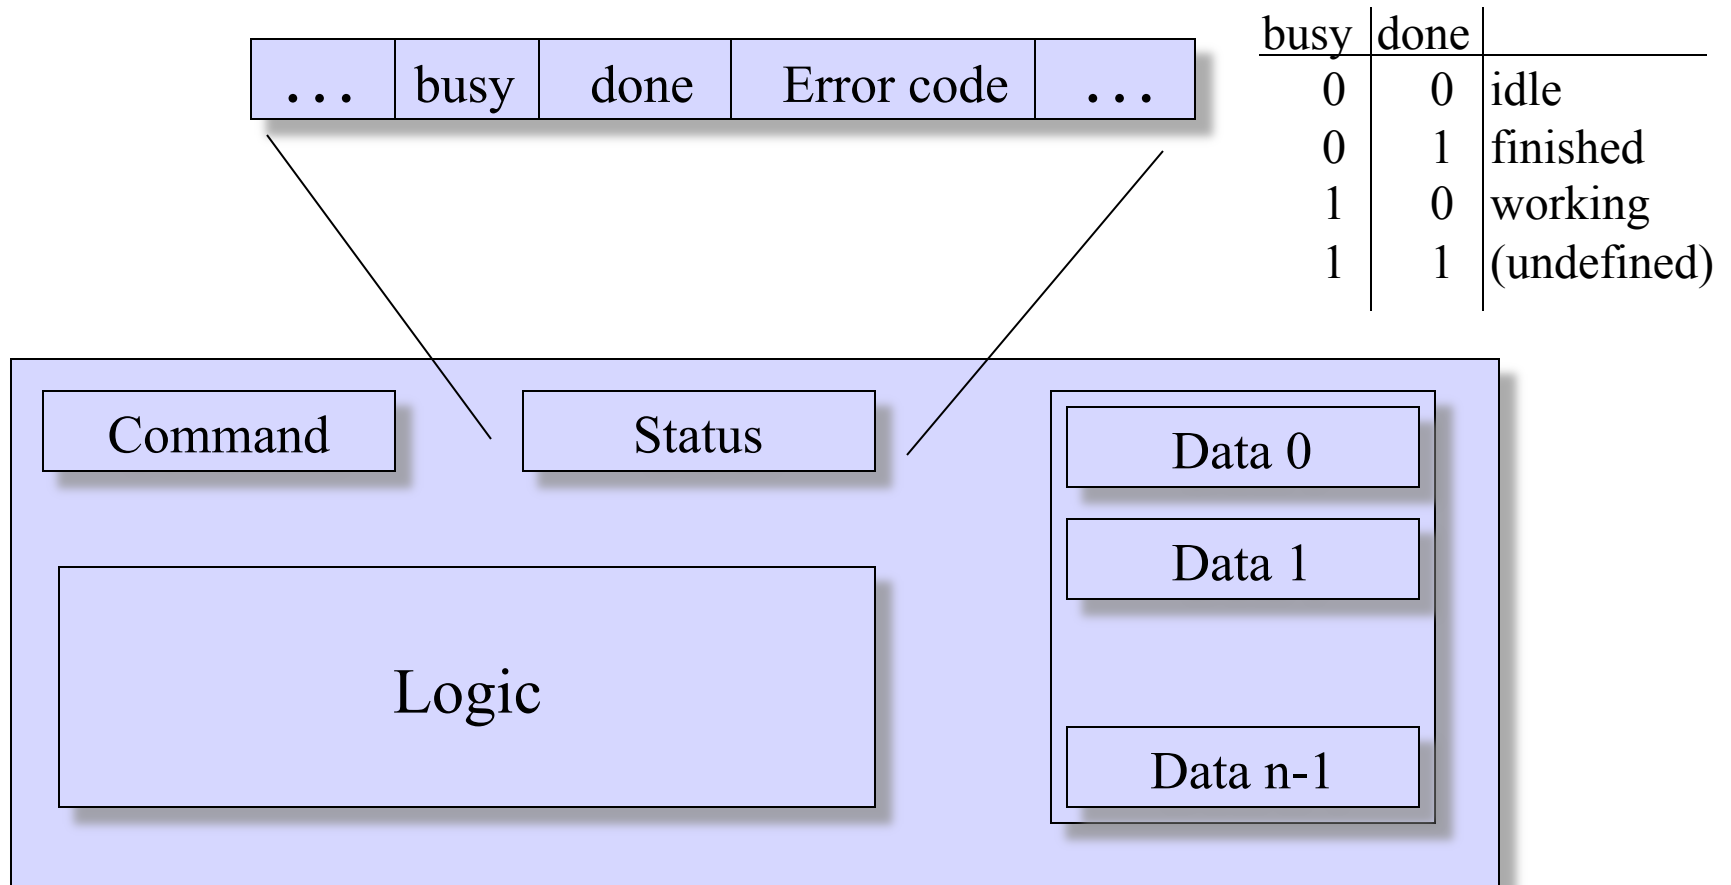


- Therefore, need 2 bits for 3 states:
  - A BUSY flag and a DONE flag
  - BUSY=0, DONE=0 => Idle
  - BUSY=1, DONE=0 => Working
  - BUSY=0, DONE=1 => Finished
  - BUSY=1, DONE=1 => Undefined

- Need three states to distinguish the following:

- Idle: no app is accessing the device
- Working: one app only is accessing the device
- Finished: the results are ready for that one app

# Device Controller Interface

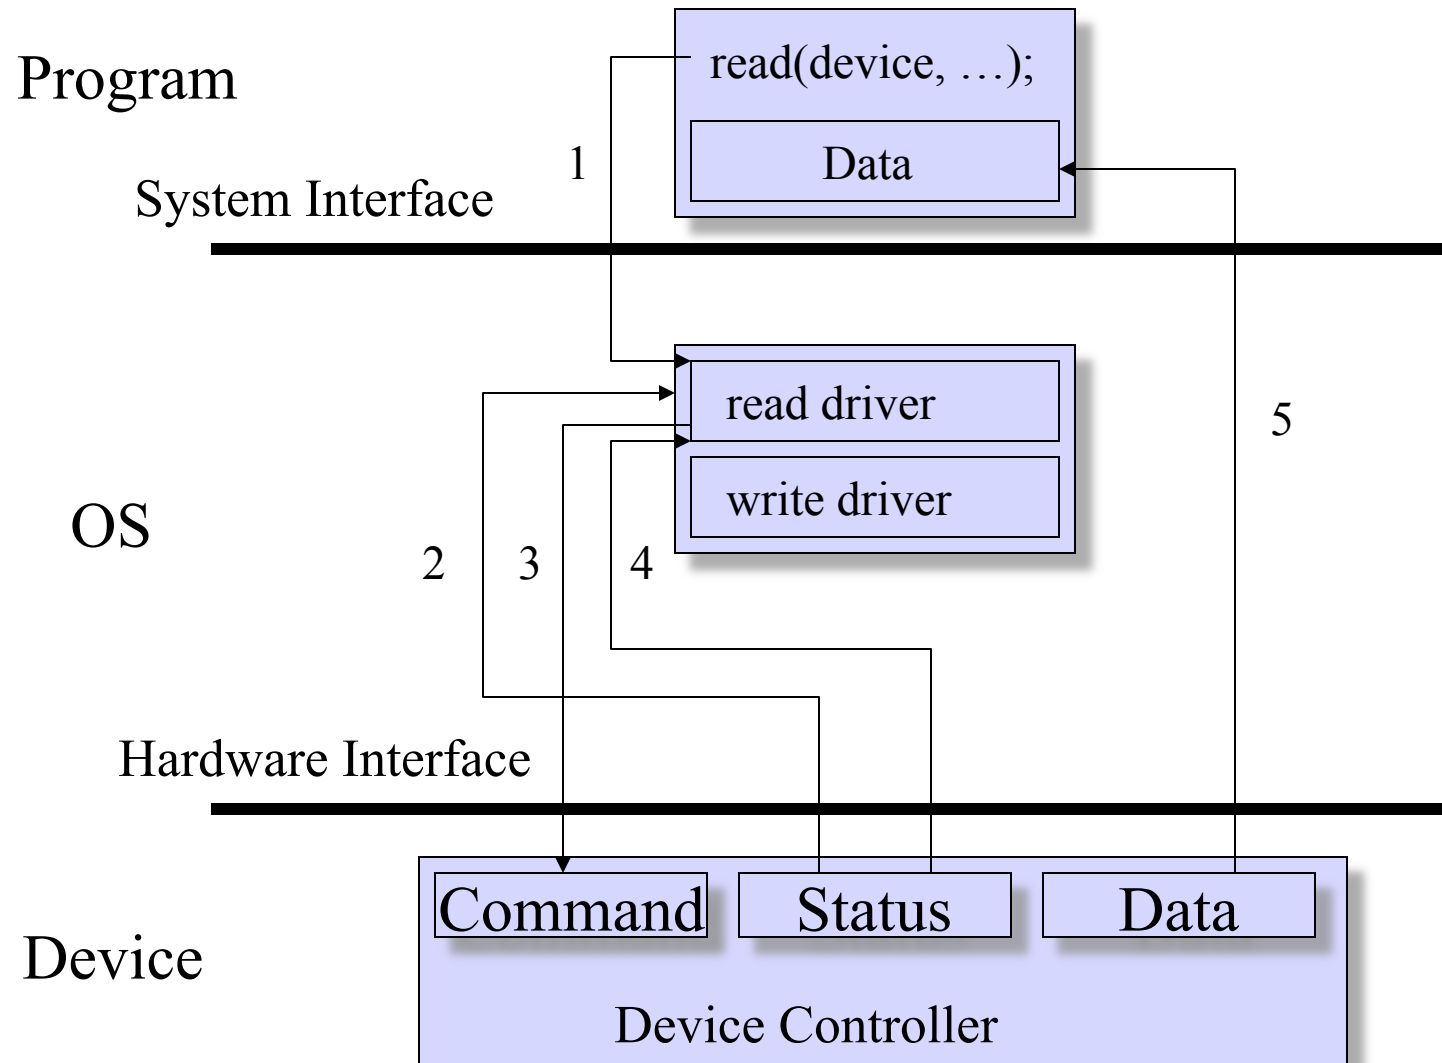


# Polling I/O: A Write Example

	BUSY	DONE
	*	*
while(deviceN.busy    deviceN.done) <waiting>;	0	0
deviceN.data[0] = <value to write>		
deviceN.command = WRITE;		
while(deviceN.busy) <waiting>;	1	0
/* finished, so read some status bits... */	0	1
deviceN.done = FALSE;	0	0

- Devices much slower than CPU
- CPU waits while device operates
- Would like to multiplex CPU to a different process while I/O is in process

# Polling I/O Read Operation





# Polling I/O – Busy Waiting

- Note that the OS is spinning in a loop twice:
  - Checking for the device to become idle
  - Checking for the device to finish the I/O request, so the results can be retrieved
  - This wastes CPU cycles that could be devoted to executing applications
- Instead, want to *overlap* CPU and I/O
  - Free up the CPU while the I/O device is processing a read/write

# Device Manager I/O Strategies

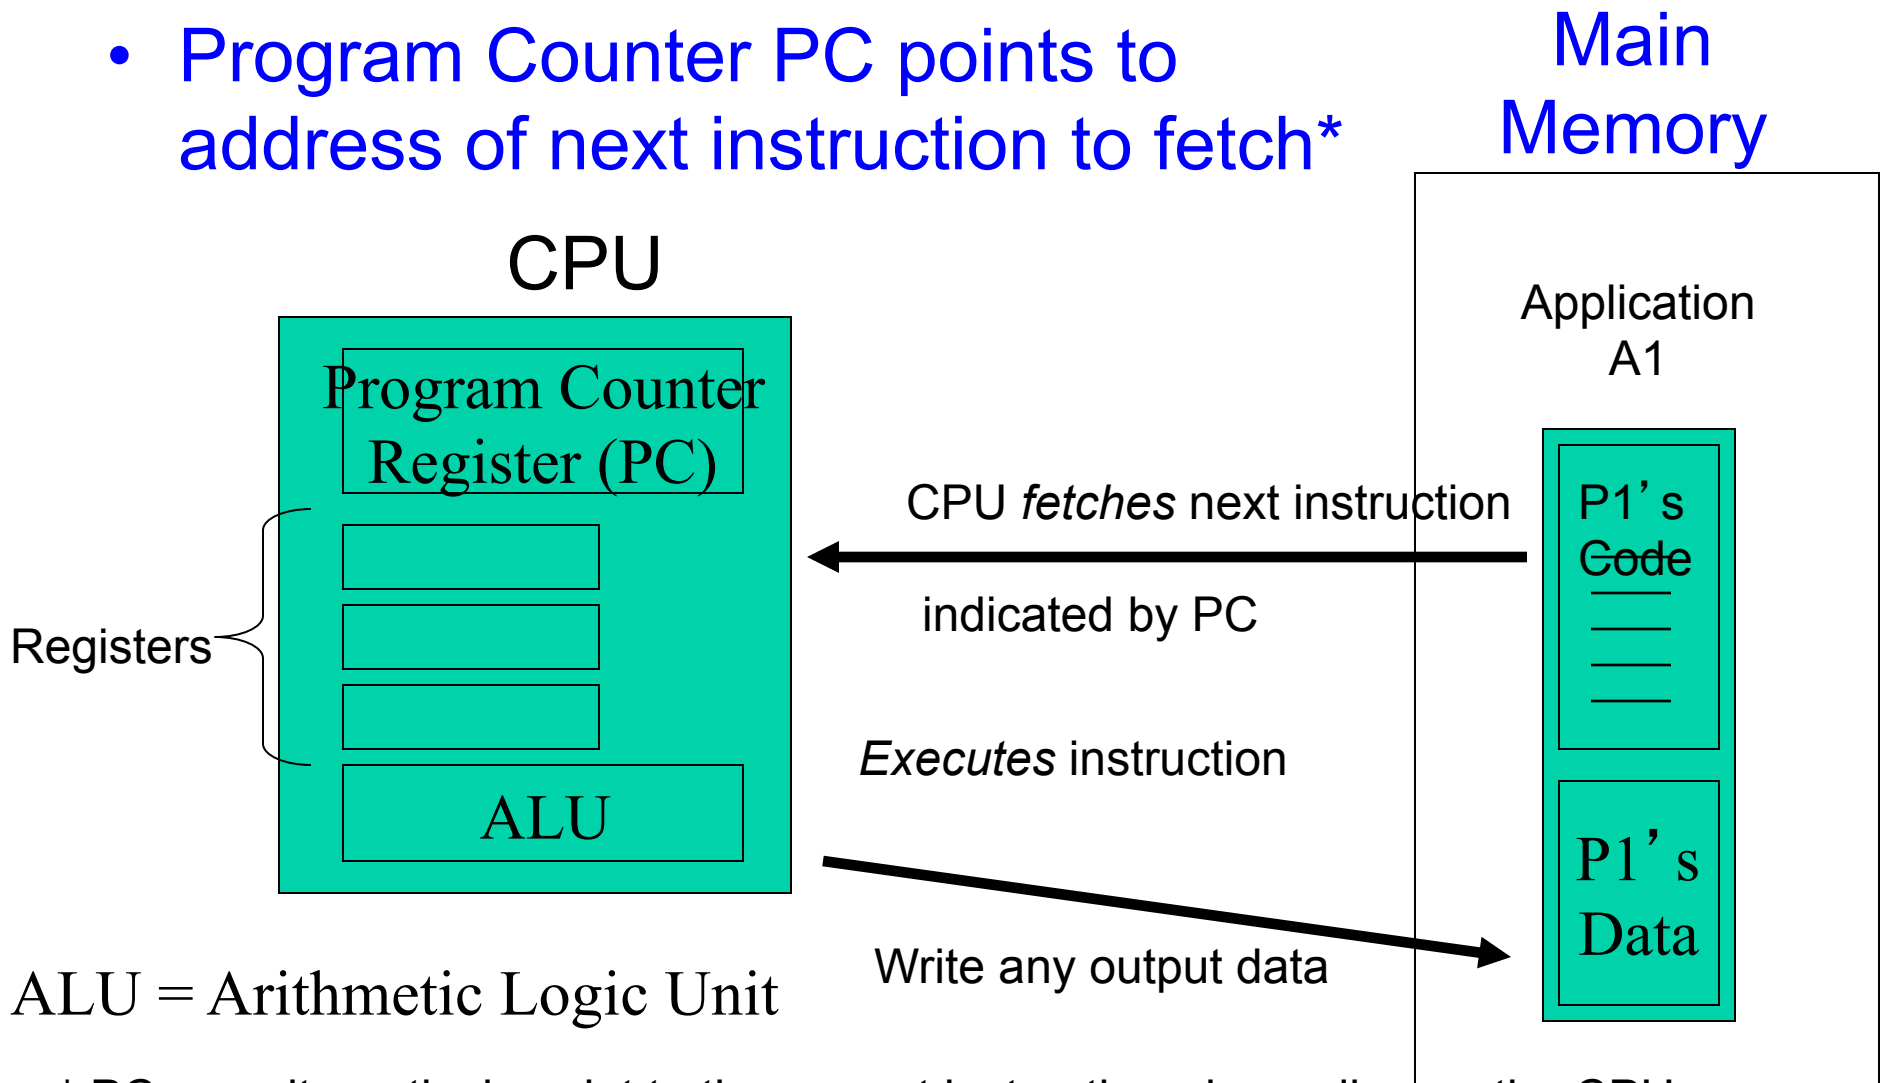
- Underneath the blocking/non-blocking synchronous/asynchronous system call API, OS can implement several strategies for I/O with devices
  - direct I/O with polling
    - the OS device manager busy-waits, we've already seen this
  - direct I/O with *interrupts*
    - More efficient than busy waiting
  - DMA with interrupts

# Hardware Interrupts

- CPU incorporates a *hardware interrupt flag*
- Whenever a device is finished with a read/write, it communicates to the CPU and raises the flag
  - Frees up CPU to execute other tasks without having to keep polling devices
- Upon an interrupt, the CPU interrupts normal execution, and invokes the OS's *interrupt handler*
  - Eventually, after the interrupt is handled and the I/O results processed, the OS resumes normal execution

# CPU Execution of a Program

- Program Counter PC points to address of next instruction to fetch\*



ALU = Arithmetic Logic Unit

\* PC can alternatively point to the current instruction, depending on the CPU

# CPU Checks Interrupt Flag Every Fetch/Execute Cycle

## CPU Pseudocode

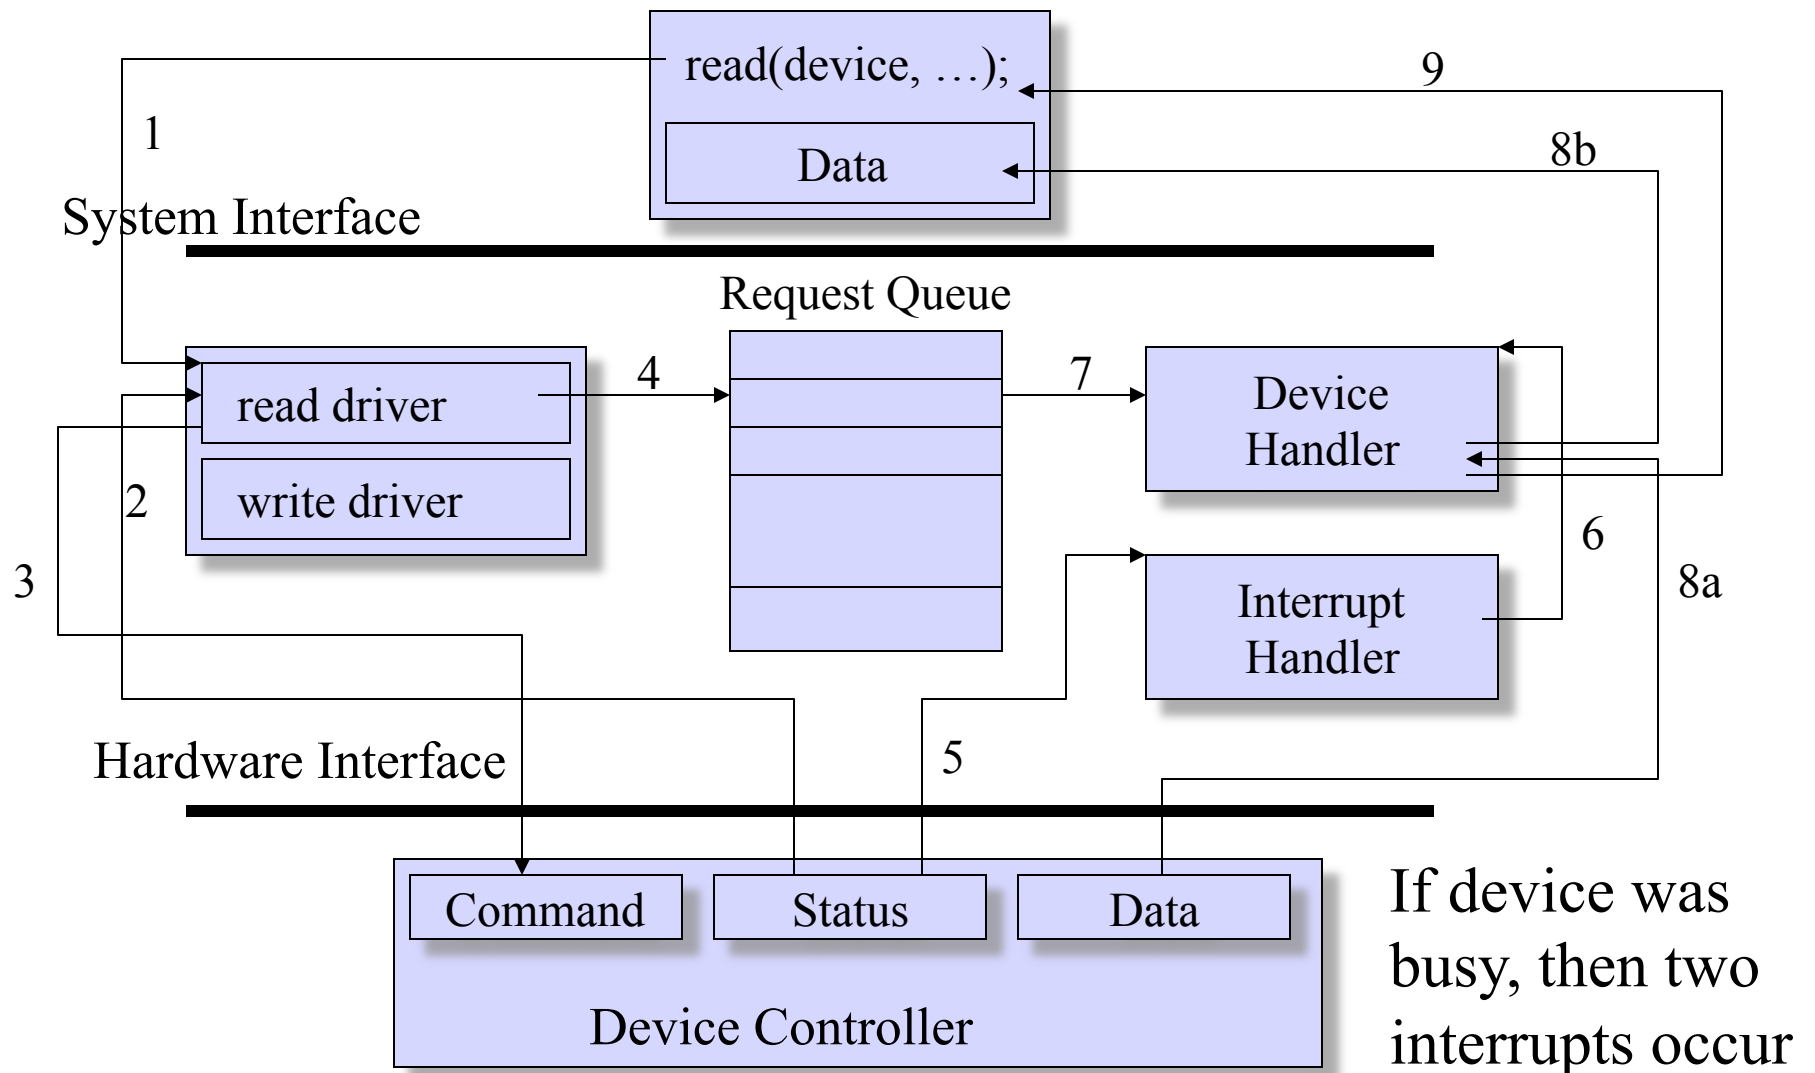
- While (no hardware failure)
  - Fetch next instruction, put in instruction register
  - Execute instruction
  - Check for interrupt: If interrupt flag enabled,
    - Save PC\*
    - Jump to interrupt handler

\* insight from Nutt's text

# Interrupt Handler

- First, save the processor state
  - Save the executing app's program counter (PC) and CPU register data
- Next, find the device causing the interrupt
  - Consult interrupt controller to find the interrupt offset, or poll the devices
- Then, jump to the appropriate device handler
  - Index into the Interrupt Vector using the interrupt offset
  - An Interrupt Service Routine (ISR) either refers to the interrupt handler, or the device handler
- Finally, reenables interrupts

# Interrupt-Driven I/O Operation

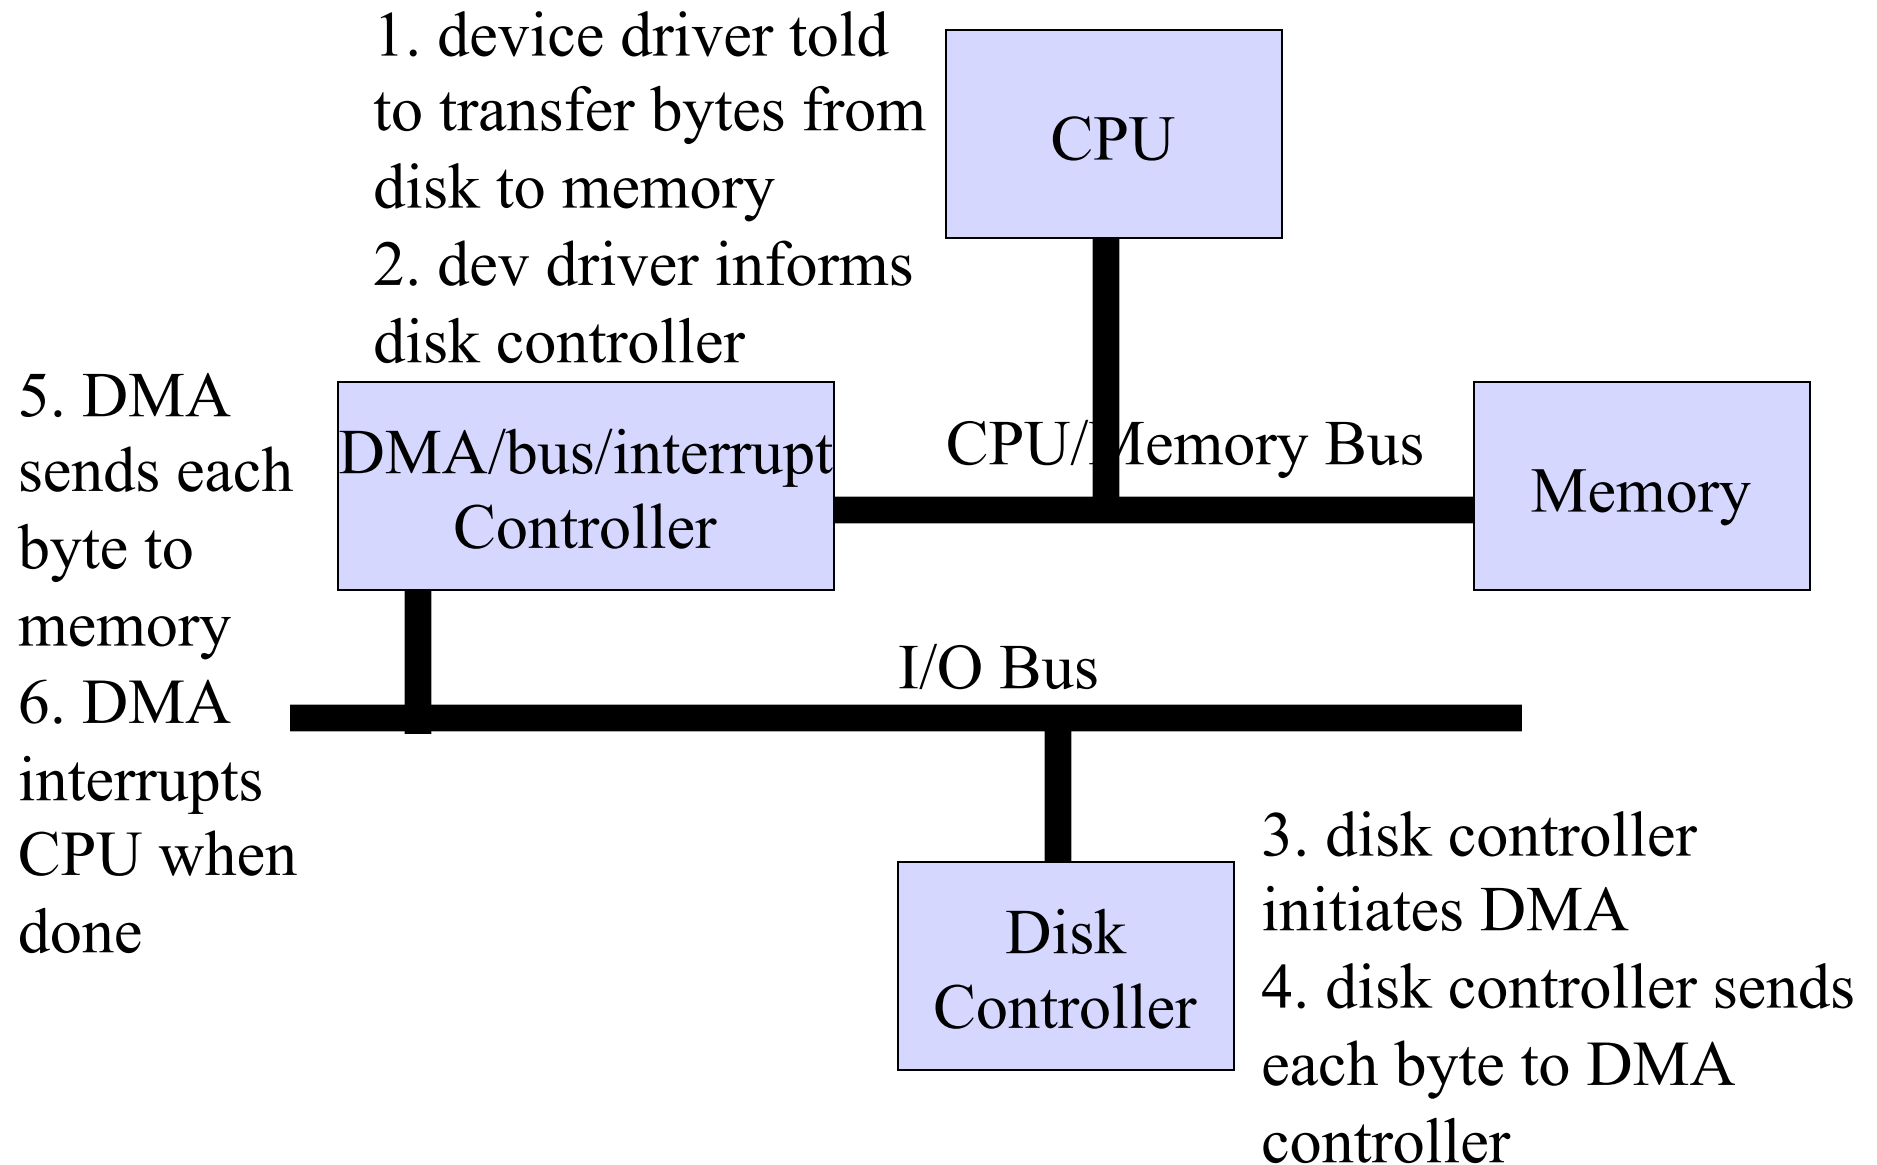


# Direct Memory Access (DMA)

- The CPU can become a bottleneck if there is a lot of I/O copying data back and forth between memory and devices
  - Example: want to copy a 1 MB file from disk into memory. The disk is only capable of delivering memory in say 1 KB blocks. So every time a 1 KB block is ready to be copied, an interrupt is raised, interrupting the CPU. This will slow down execution of normal programs and the OS.
  - Worst cases: CPU could be interrupted after the transfer of every byte/character, or every packet from the network card
- DMA solution: Bypass the CPU for large data copies, and only raise an interrupt at the very end of the data transfer, instead of at every intermediate block



# DMA with Interrupts Example



# Direct Memory Access (DMA)

- Since both CPU and the DMA controller have to move data to/from main memory, how do they share main memory?
  - Burst mode
    - While DMA is transferring, CPU is blocked from accessing memory
  - Interleaved mode or “cycle stealing”
    - DMA transfers one word to/from memory, then CPU accesses memory, then DMA, then CPU, etc... - interleaved
  - Transparent mode – DMA only transfers when CPU is not using the system bus
    - Most efficient but difficult to detect

# Memory-Mapped I/O

- Non-memory mapped (port or port-mapped) I/O typically requires special I/O machine instructions to read/write from/to device controller registers
  - e.g. on Intel x86 CPUs, have IN, OUT
    - Example: OUT dest, src (using Intel syntax, not Gnu syntax)
      - Writes to a device port dest from CPU register src
    - Example: IN dest, src
      - Reads from a device port src to CPU register src
    - Only OS in kernel mode can execute these instructions
    - Later Intel introduced INS, OUTS (for strings), and INSB/INSW/INSD (different word widths), etc.

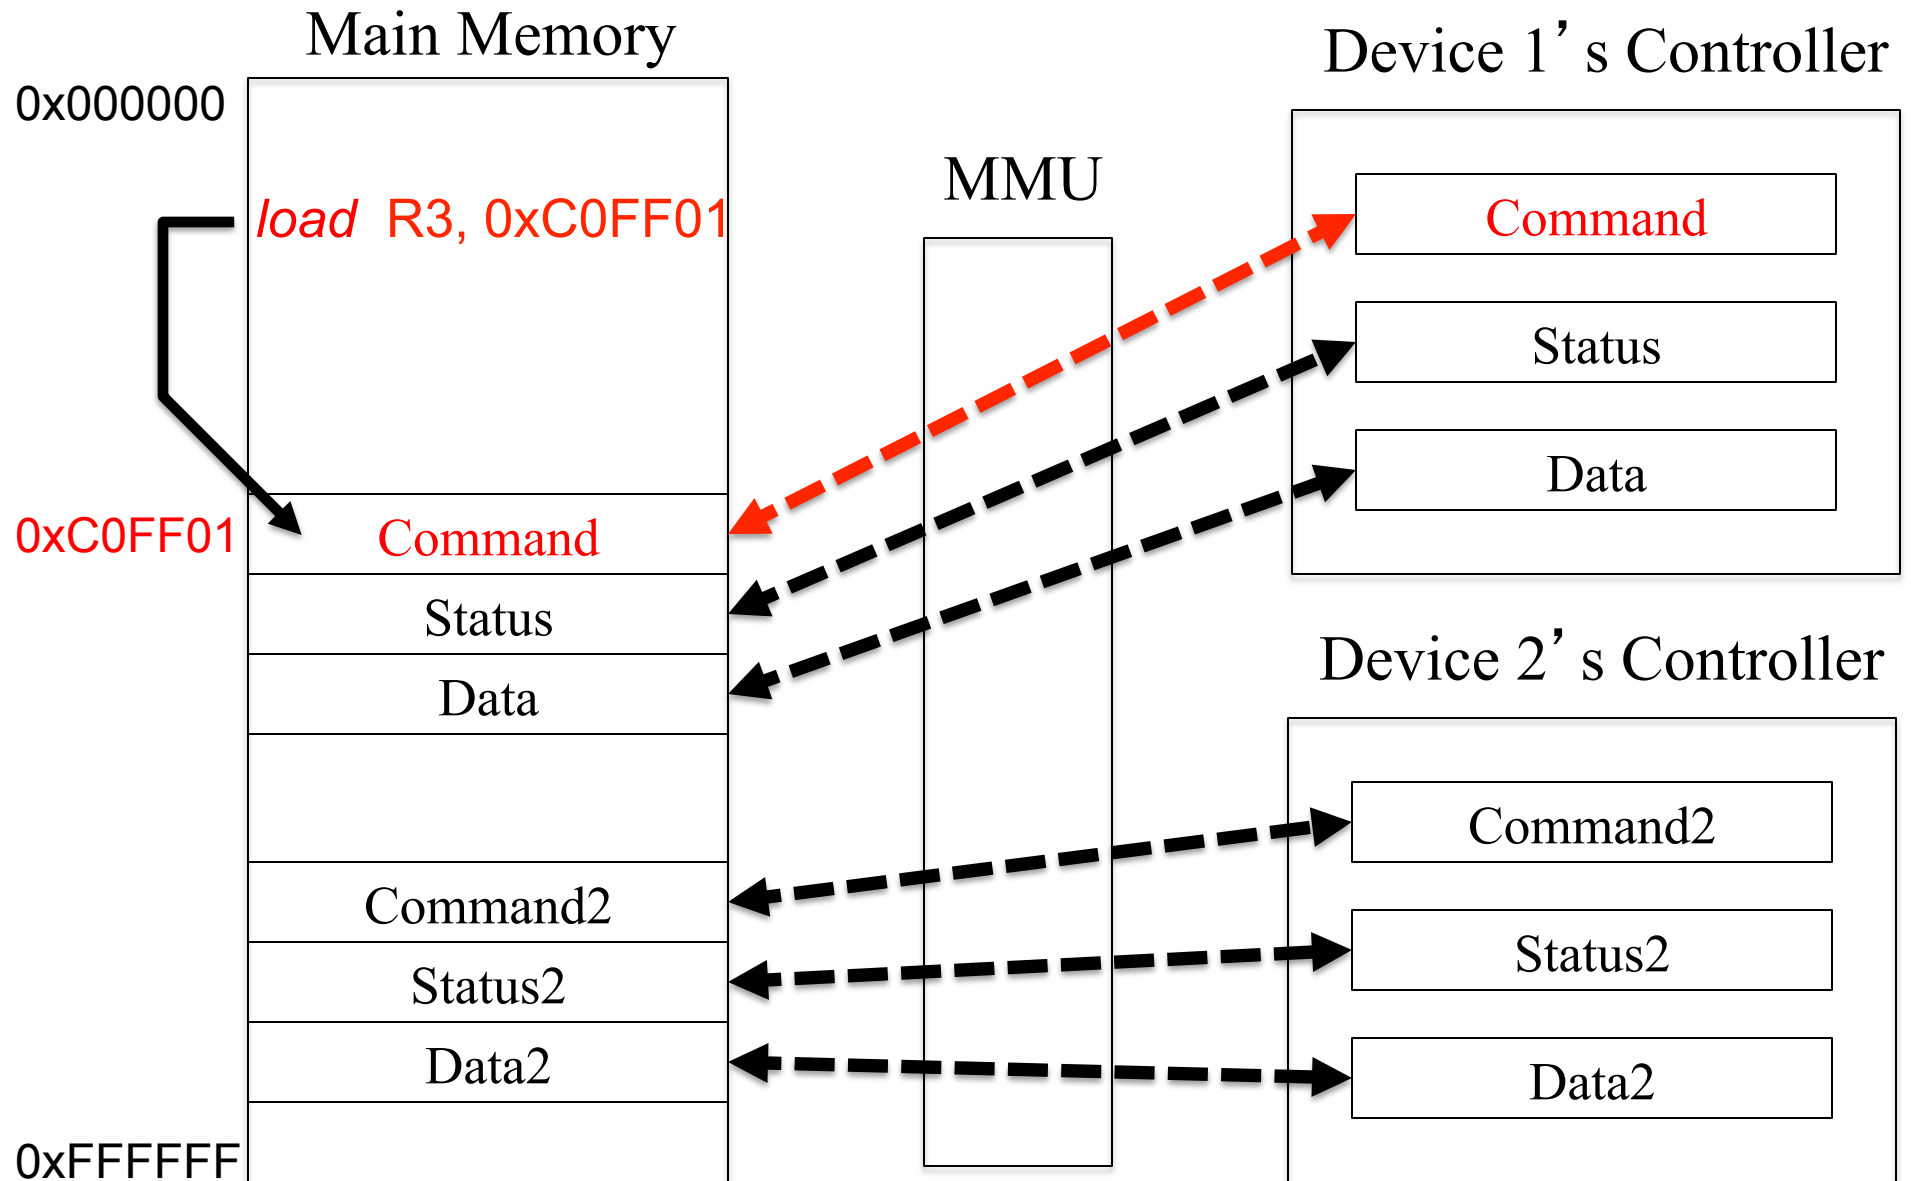
# Memory-Mapped I/O (2)

- port-mapped I/O is quite limited
  - IN and OUT can only store and load
  - don't have full range of memory operations for normal CPU instructions
    - Example: to increment the value in say a device's data register, have to copy register value into memory, add one, and copy it back to device register.
      - With memory-mapped I/O, can increment value in memory directly, and it gets reflected into the device controller's data register automatically
    - AMD did not extend the port I/O instructions when defining the x86-64

# Memory-Mapped I/O (3)

- With memory-mapped I/O, just address memory directly using normal instructions to speak to an I/O address
  - e.g. `load R3, 0xC0FF01`
  - the memory address 0xC0FF01 is mapped to the I/O device's registers
  - Memory Management Unit (MMU) maps memory values and data to/from device registers
    - Device registers are assigned to a block of memory
    - When a value is written into that I/O-mapped memory, the device sees the value, loads the appropriate value and executes the appropriate command

# Memory-Mapped I/O (4)



# Memory-Mapped I/O (5)

- Typically, devices are mapped into lower memory
  - frame buffers for displays take the most memory, since most other devices have smaller buffers
  - Even a large display might take only 10 MB of memory, which in modern address spaces of tens-hundreds of GBs is quite modest – so memory-mapped I/O is a small penalty

## Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



# Device Independent Part

- A set of system calls that an application program can use to invoke I/O operations
- A particular device will respond to only a subset of these system calls
  - A keyboard does not respond to *write()* system call
- POSIX set: *open()*, *close()*, *read()*, *write()*, *lseek()* and *ioctl()*.

# Device Independent Function Call

Trap Table

func<sub>i</sub>(...)

```
dev_func_i(devID, ...) {  
    // Processing common to all devices  
    ...  
    switch(devID) {  
        case dev0: dev0_func_i(...);  
                    break;  
        case dev1: dev1_func_i(...);  
                    break;  
        ...  
        case devM: devM_func_i(...);  
                    break;  
    };  
    // Processing common to all devices  
    ...  
}
```

# Adding a New Device

- Write device-specific functions for each I/O system call
- For each system call, add a new *case* clause to the *switch* statement in device independent function call
- Compile the kernel and new drivers

Problem: Need to compile the kernel, every time a new device or a new driver is added

# Loadable Kernel Modules

- A loadable kernel module (LKM) is an object file that contains code to extend a running kernel
- Windows, Linux, OS X ...
- Without loadable kernel modules, an OS would have to include all possible anticipated functionality already compiled directly into the base kernel
- Linux (kernel object: *.ko* extension)
  - *modprobe ( )* high level handling of LKMs (add or remove)
  - *lsmod* to list all loaded LKMs
  - *Insmod ( )* to insert and LKM
  - *rmmod ( )* to remove an LKM
  - See */lib/modules* for the all the LKMs

# *insmod ( )* command

- *insmod* makes an *init\_module* system call to load the LKM into kernel memory
  - *sudo insmod helloworld.ko*
- *init\_module* system call invokes the LKM's initialization routine (also called *init\_module*) right after it loads the LKM
- The LKM author sets up the initialization routine to call a kernel function that registers the subroutines that the LKM contains

# LKM example: Reconfigurable Device Drivers

- Allows system administrators to add a device driver to the OS without recompiling the OS
- The new driver is first stored as a *.ko* file
  - Contains an initialization routine
- The initialization routine calls a kernel function to register the device
  - e.g. *register\_chrdev*, *register\_blkdev*

- An entry table stores the actual function pointers for each device specific function call

`dev_func_i[N]`

- Replace *switch* statement with

`dev_func_i[j] (...);`

- Device registration: Fill appropriate function pointers in the entry table

# Universal Serial Bus (USB)

- USB is an industry standard that defines the cables, connectors and communication protocols used in a bus for connection, communication and power supply between computers and electronic devices
  - Keyboards, mouse, printers, digital cameras, etc.
- USB has effectively replaced a variety of earlier interfaces such as serial and parallel ports as well as chargers for portable devices



# Firewire (IEEE 1394)

- A serial bus interface for high speed communication and real-time data transfer
- Comparable to USB