

# CSCI 3753

# Operating Systems

## Processes

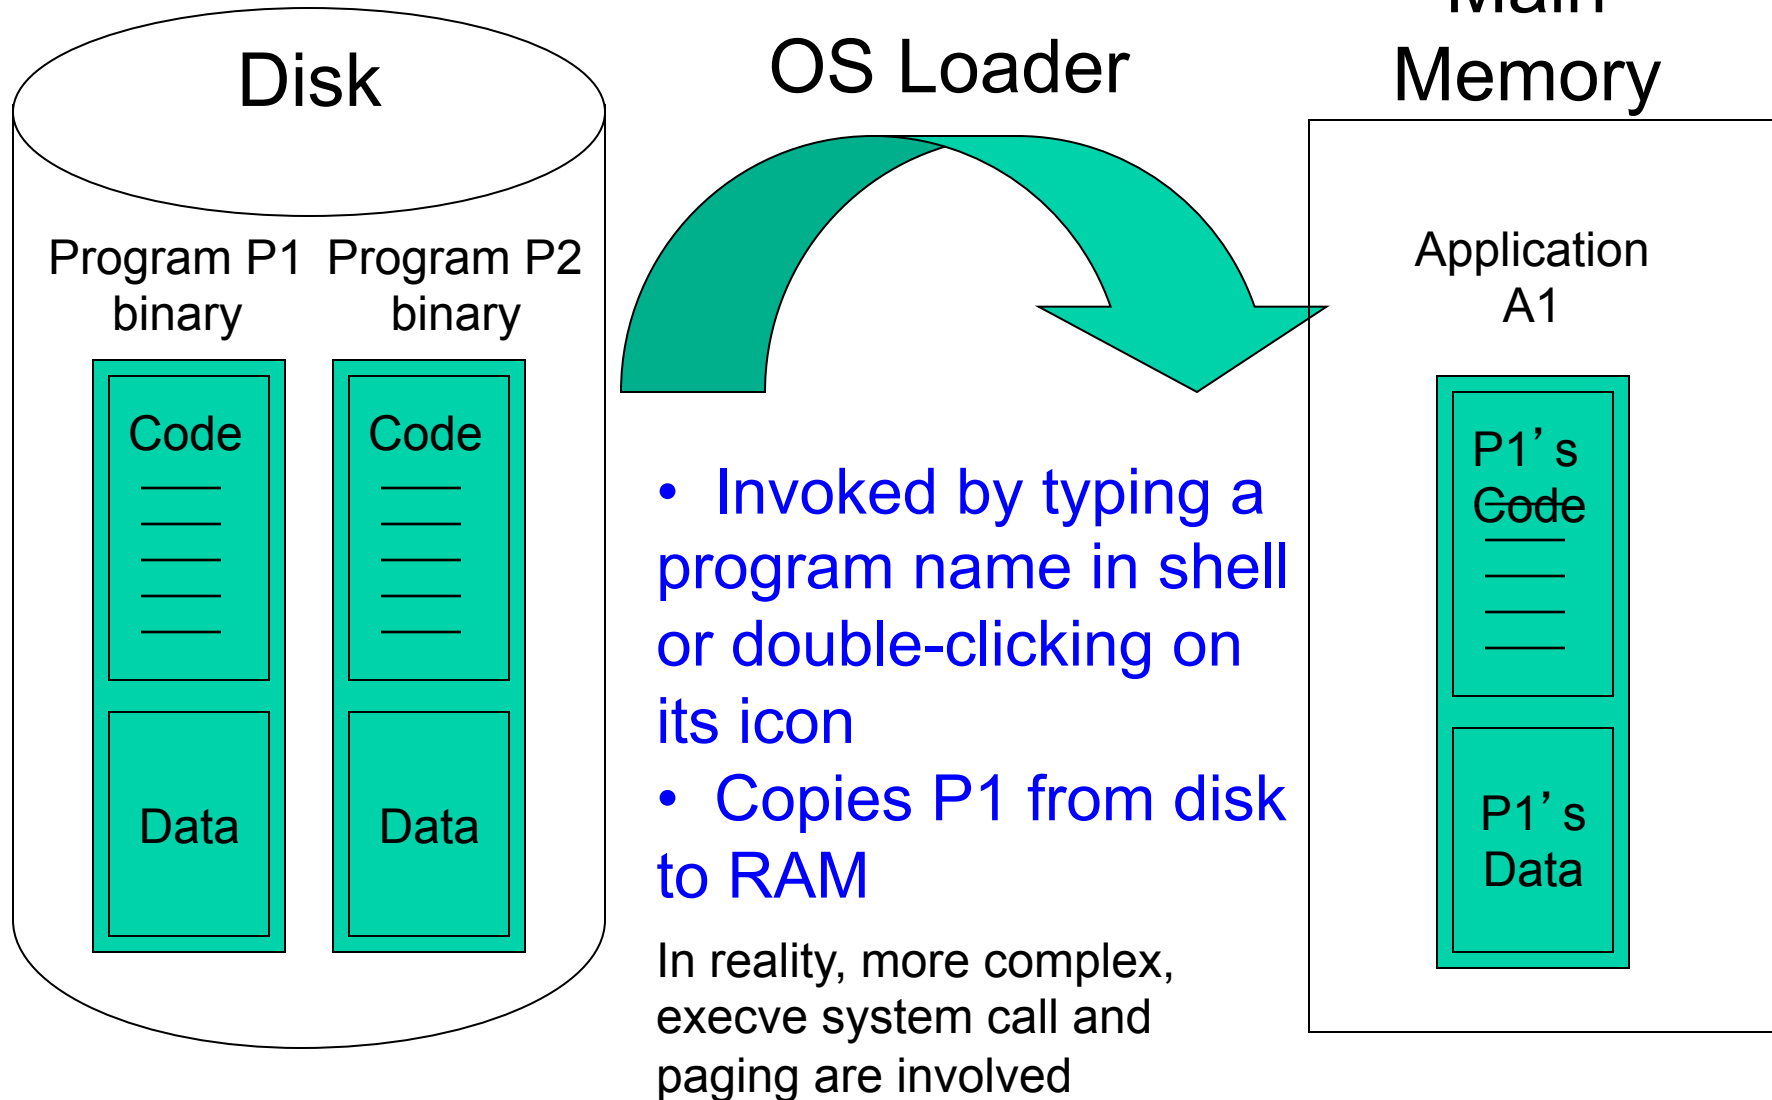
Lecture Notes By

Shivakant Mishra

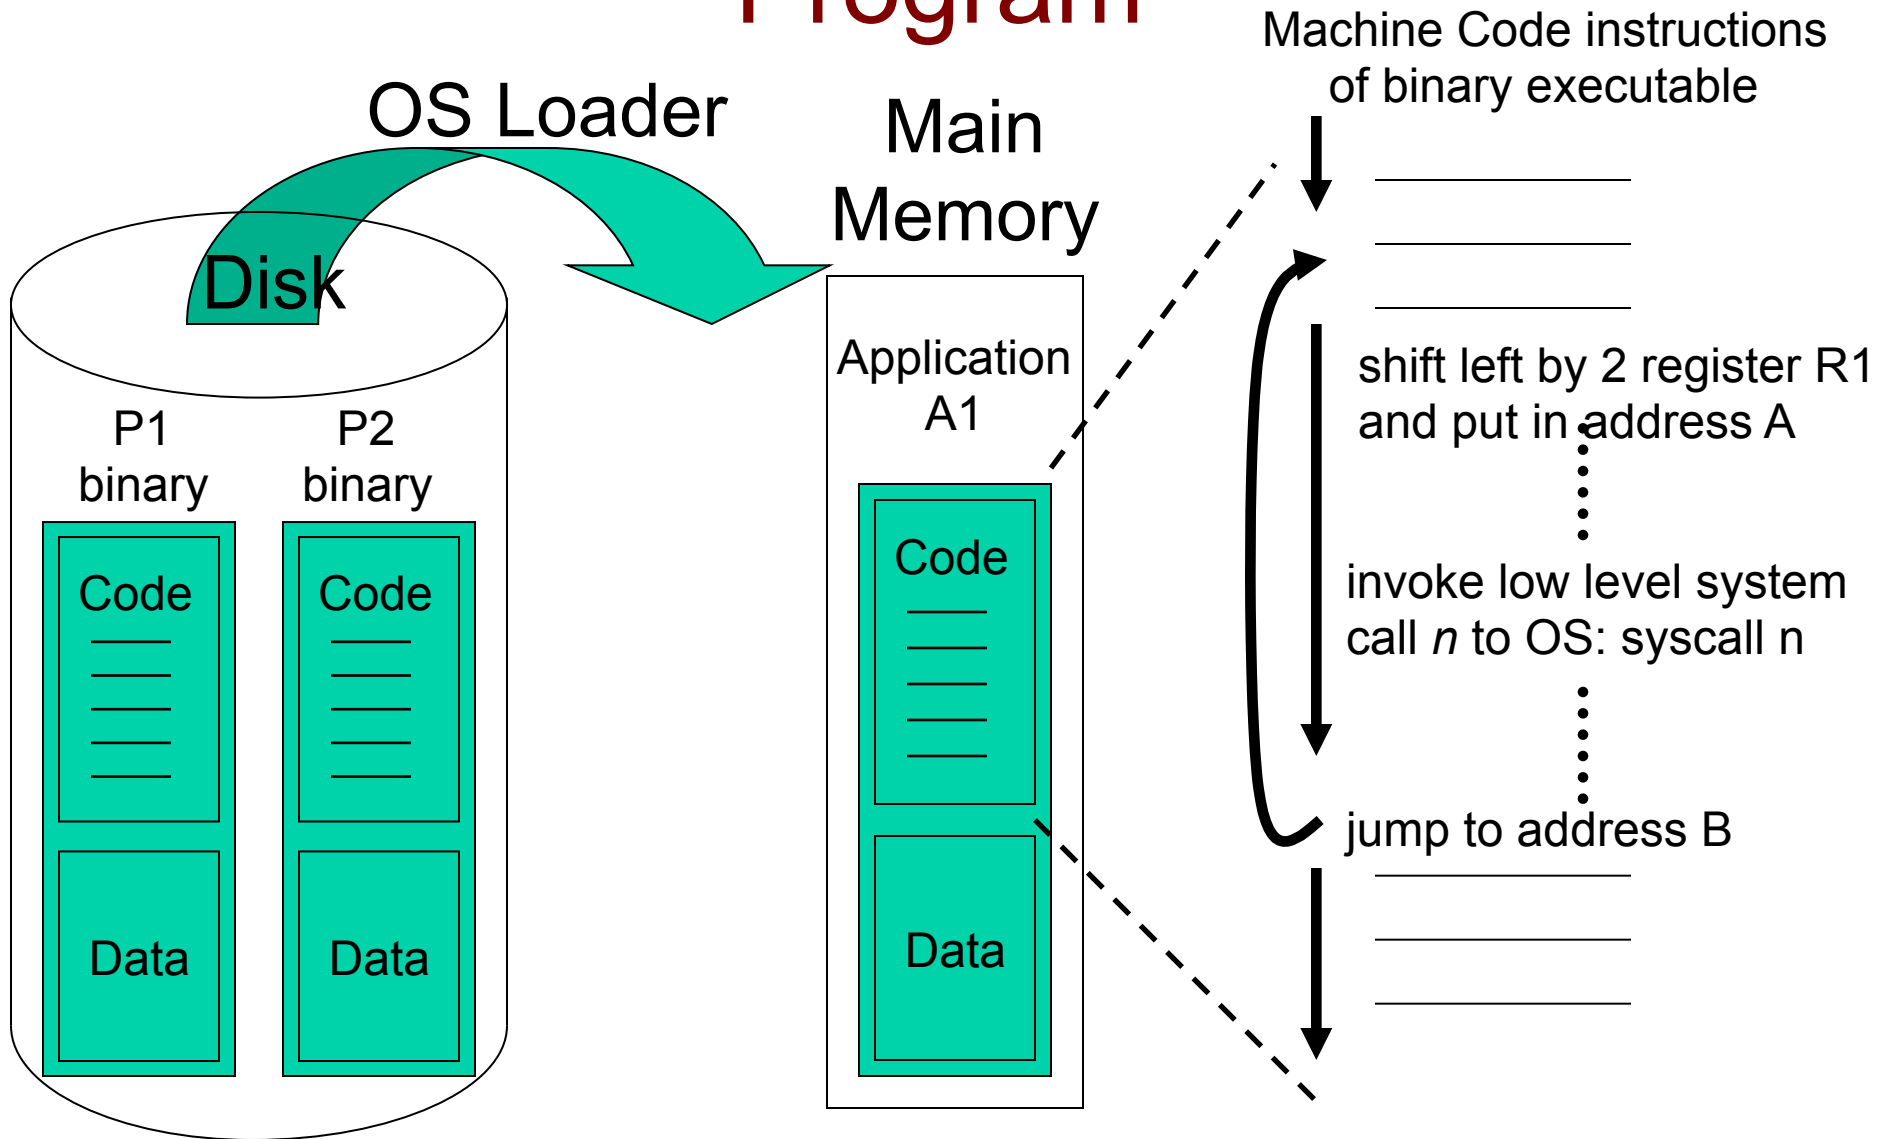
Computer Science, CU-Boulder

Last Update: 01/31/14

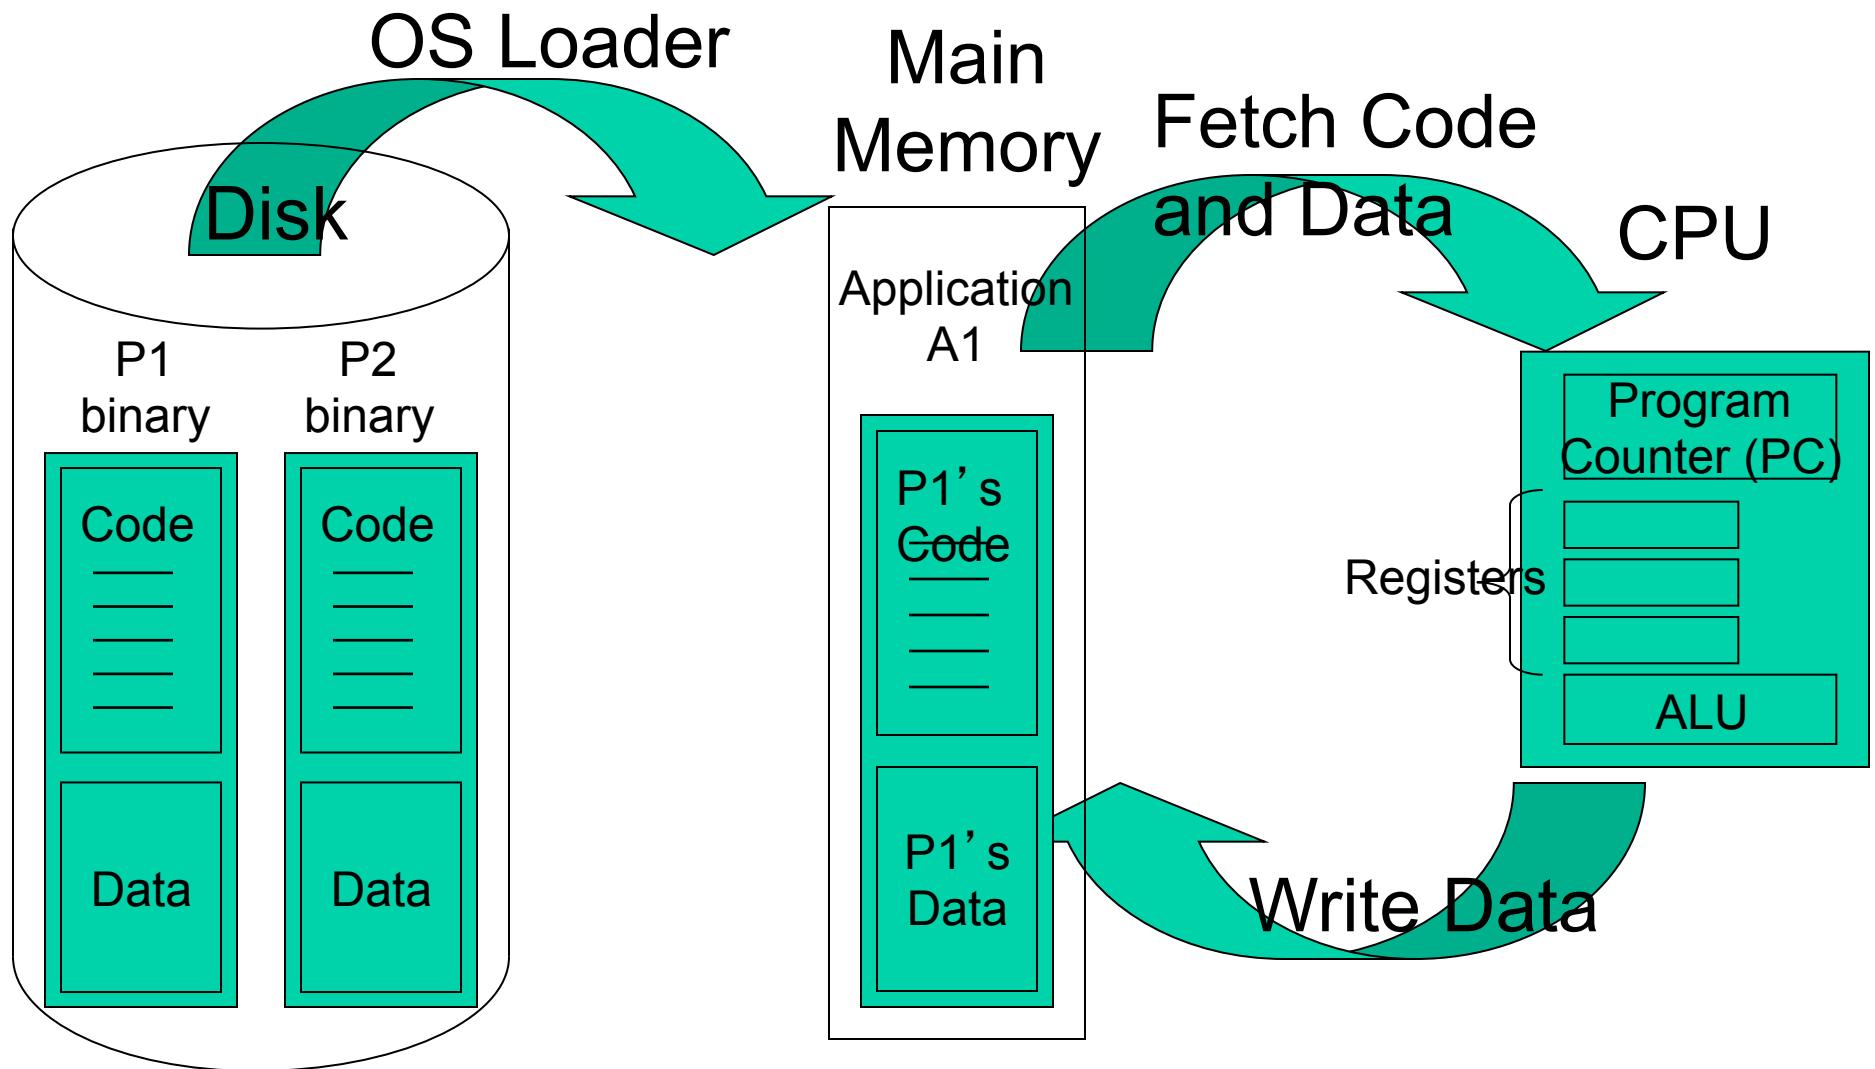
# Loading a Program into Memory



# Loading and Executing a Program

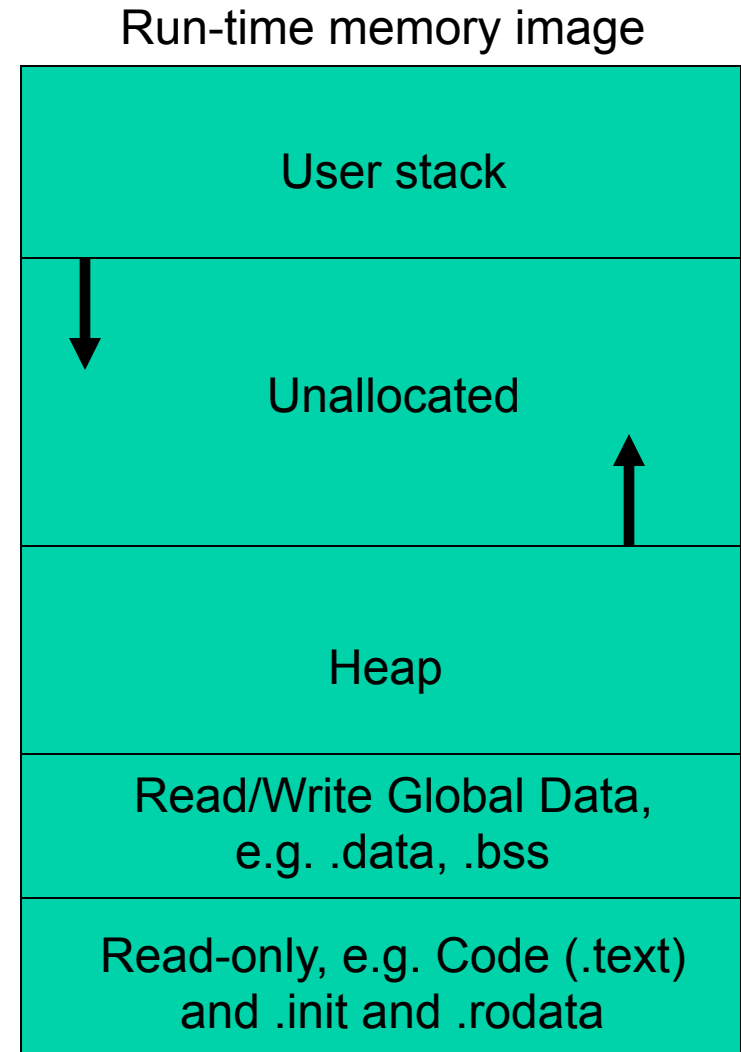


# Loading and Executing a Program



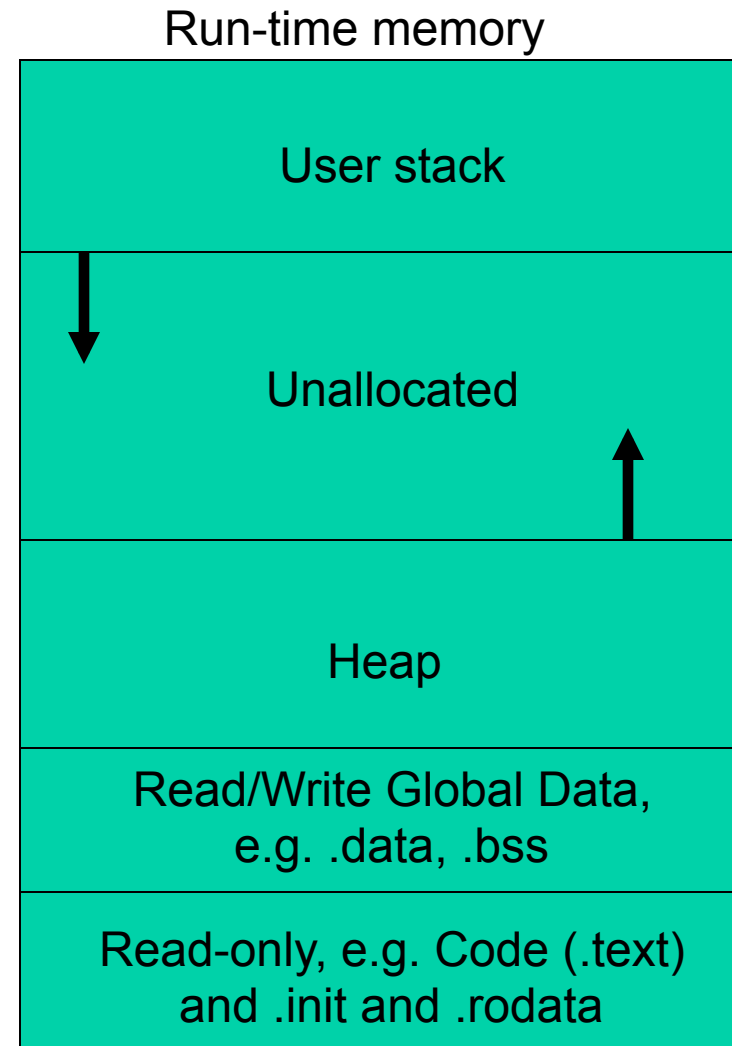
# Loading Executable Object Files

- When a program is loaded into RAM, it becomes an actively executing application
- The OS allocates a stack and heap to the app in addition to code and global data.
  - A call stack is for local variables, function parameters and return addresses
  - A heap is for dynamic variables, e.g. *malloc()*, *new*
  - Usually, stack grows downward from high memory, heap grows upward from low memory, but this architecture-specific

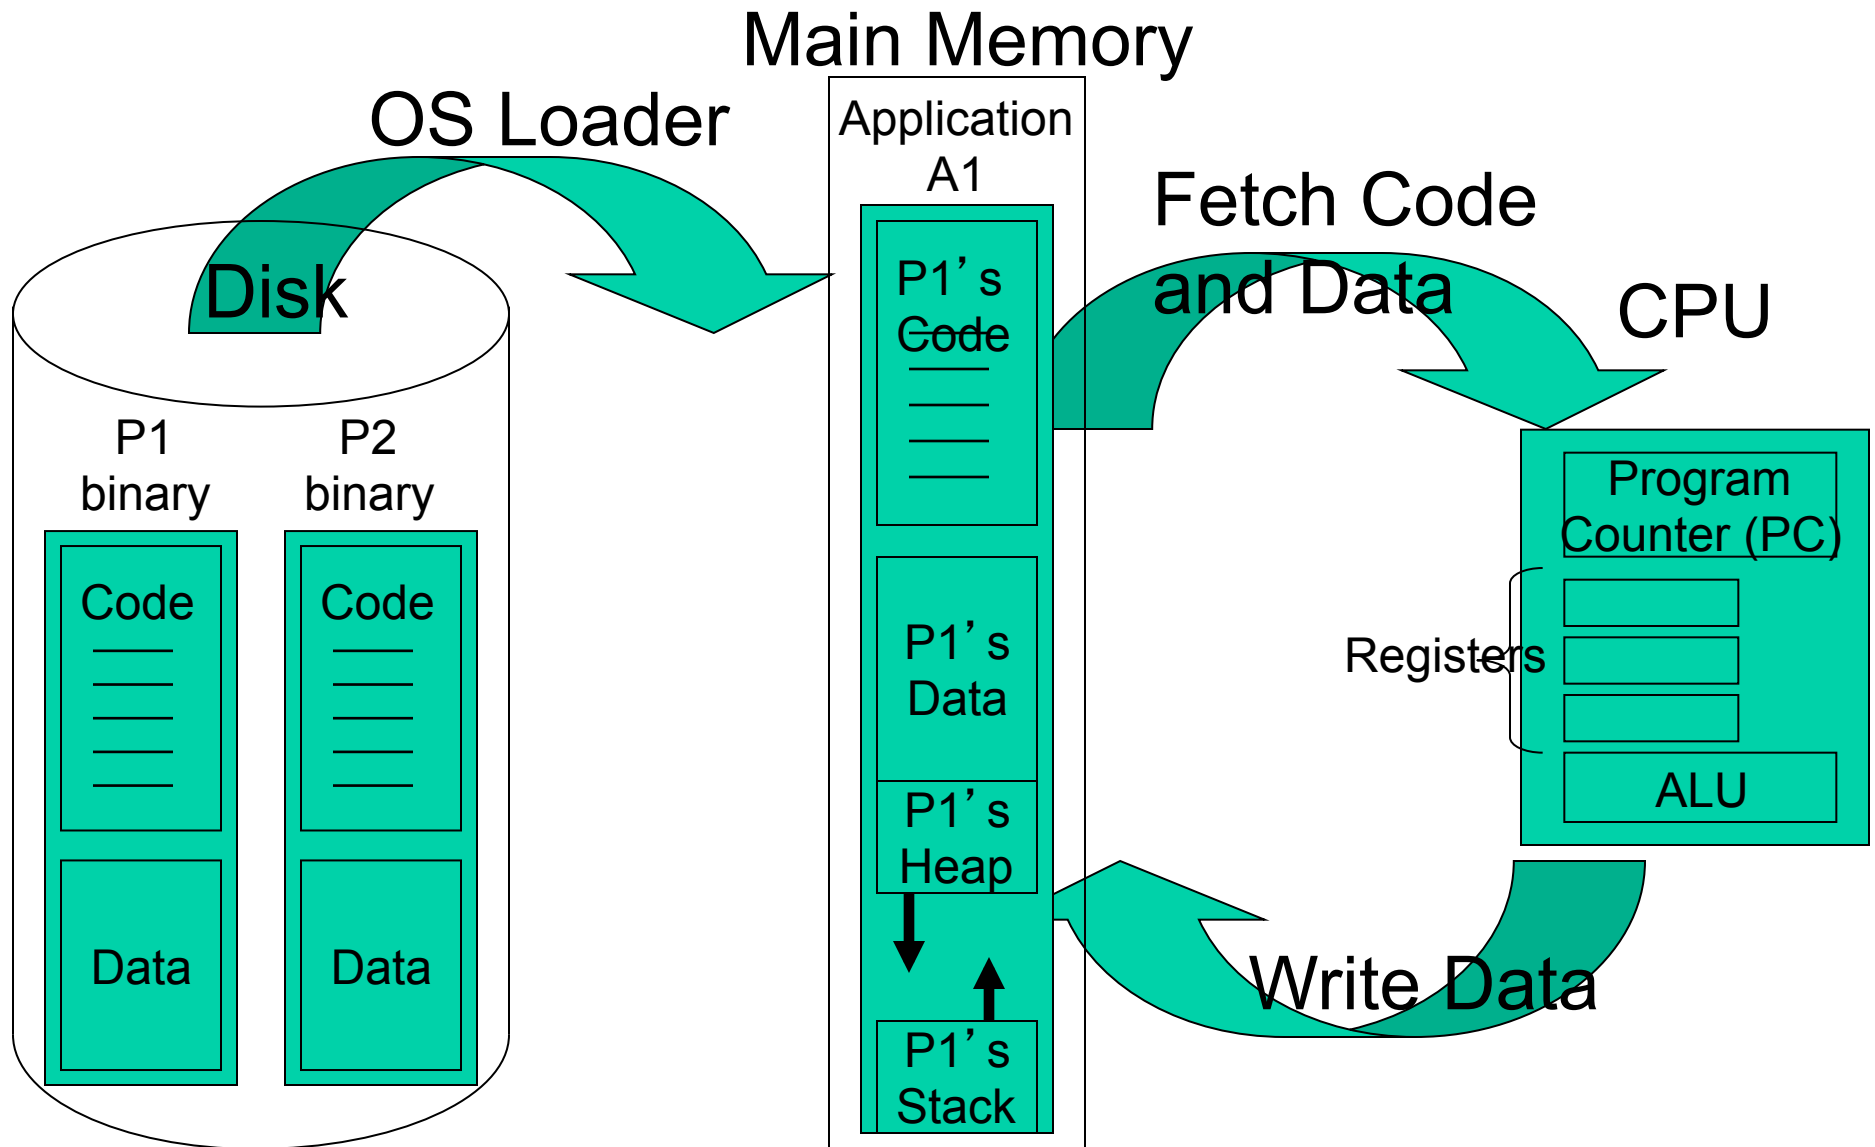


# Running Executable Object Files

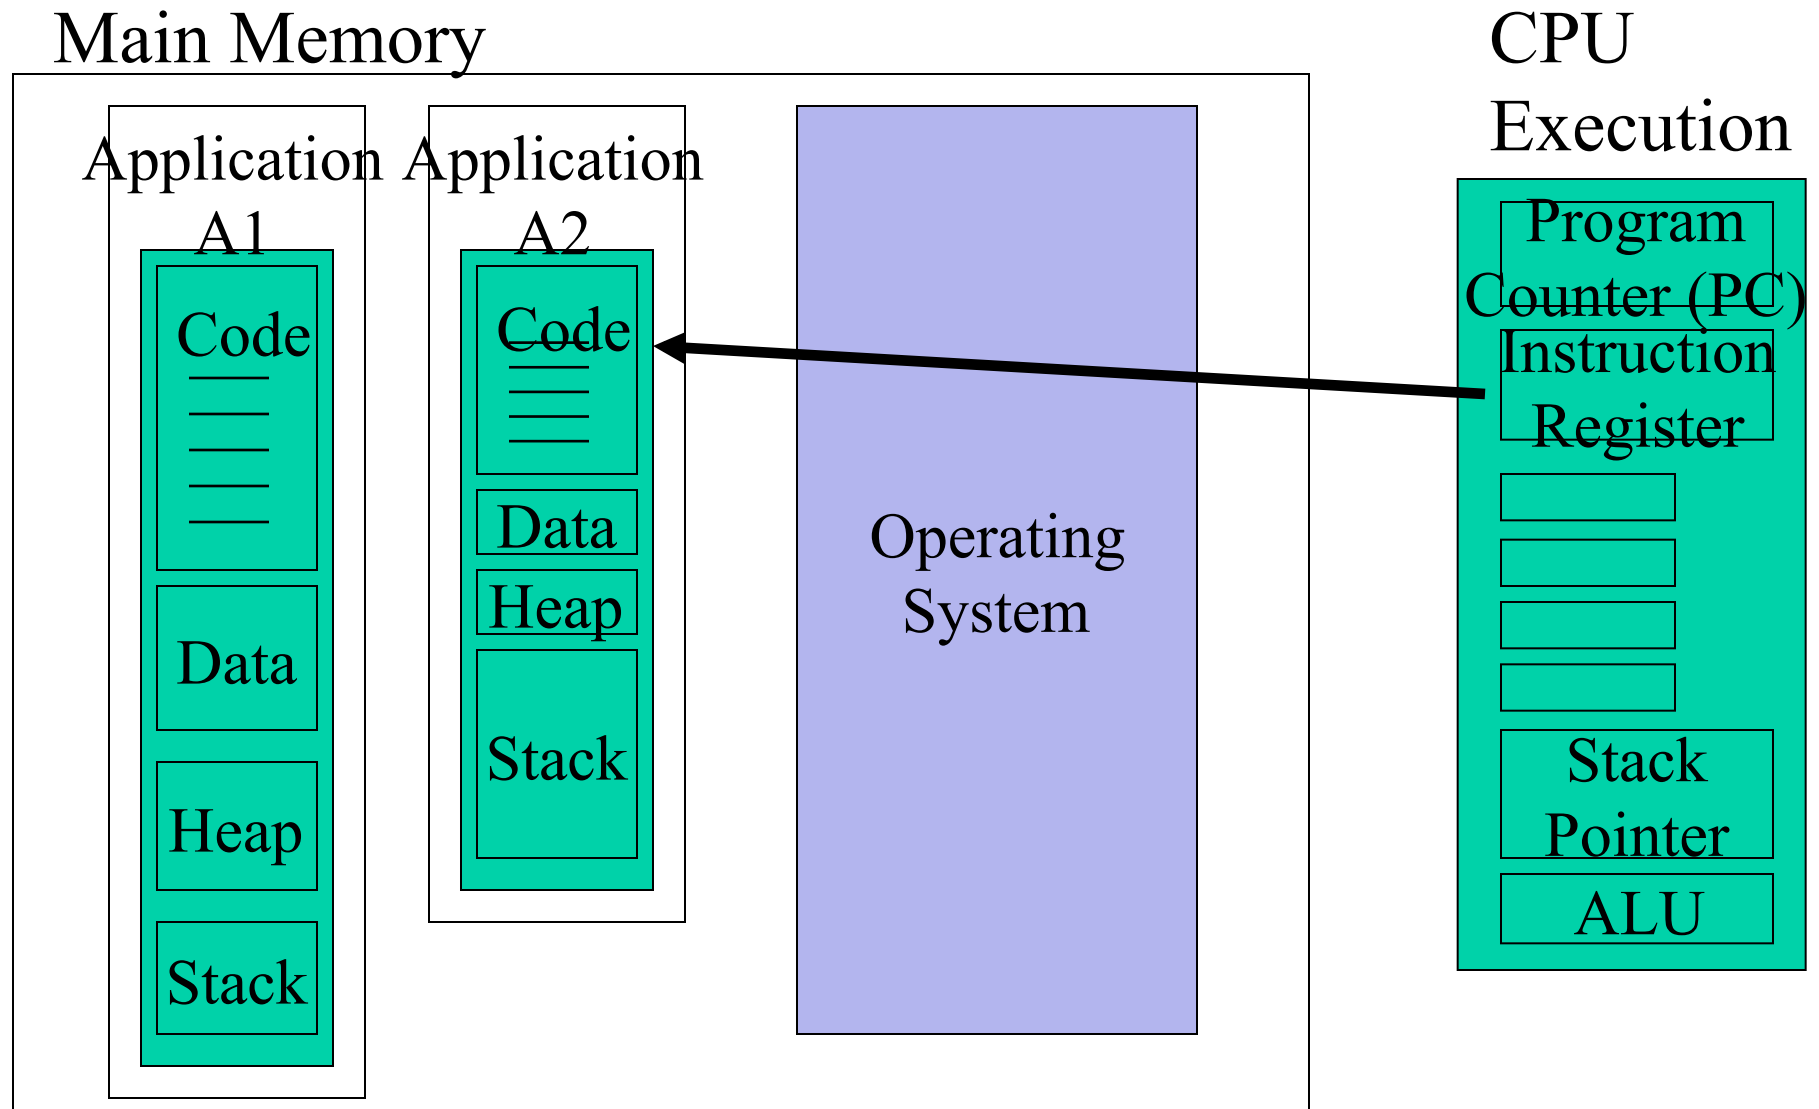
- Stack contains local variables
  - As `main()` calls function `f1`, we allocate `f1`'s local variables on the stack
  - If `f1` calls `f2`, we allocate `f2`'s variables on the stack below `f1`'s, thereby growing the stack, etc...
  - When `f2` is done, we deallocate `f2`'s local variables, popping them off the stack, and return to `f1`
- Stack dynamically expands and contracts as program runs and different levels of nested functions are called
- Heap contains run-time variables/buffers
  - Obtained from `malloc()`
  - Program should `free()` the `malloc`'ed memory
- Heap can also expand and contract during program execution



# Loading and Executing a Program – a more complete picture



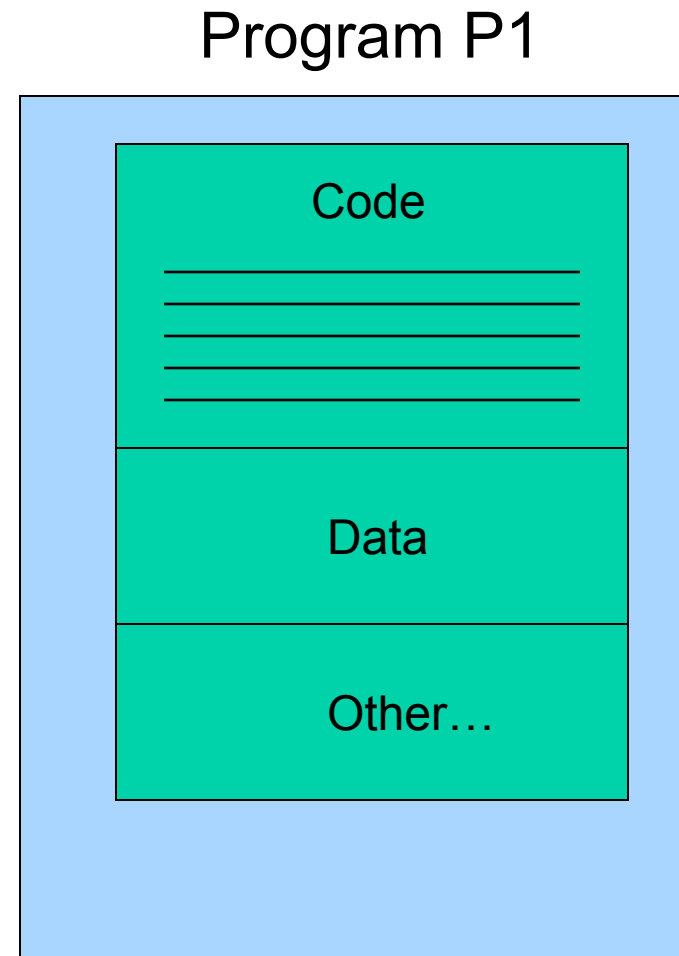
# Multiple Applications + OS



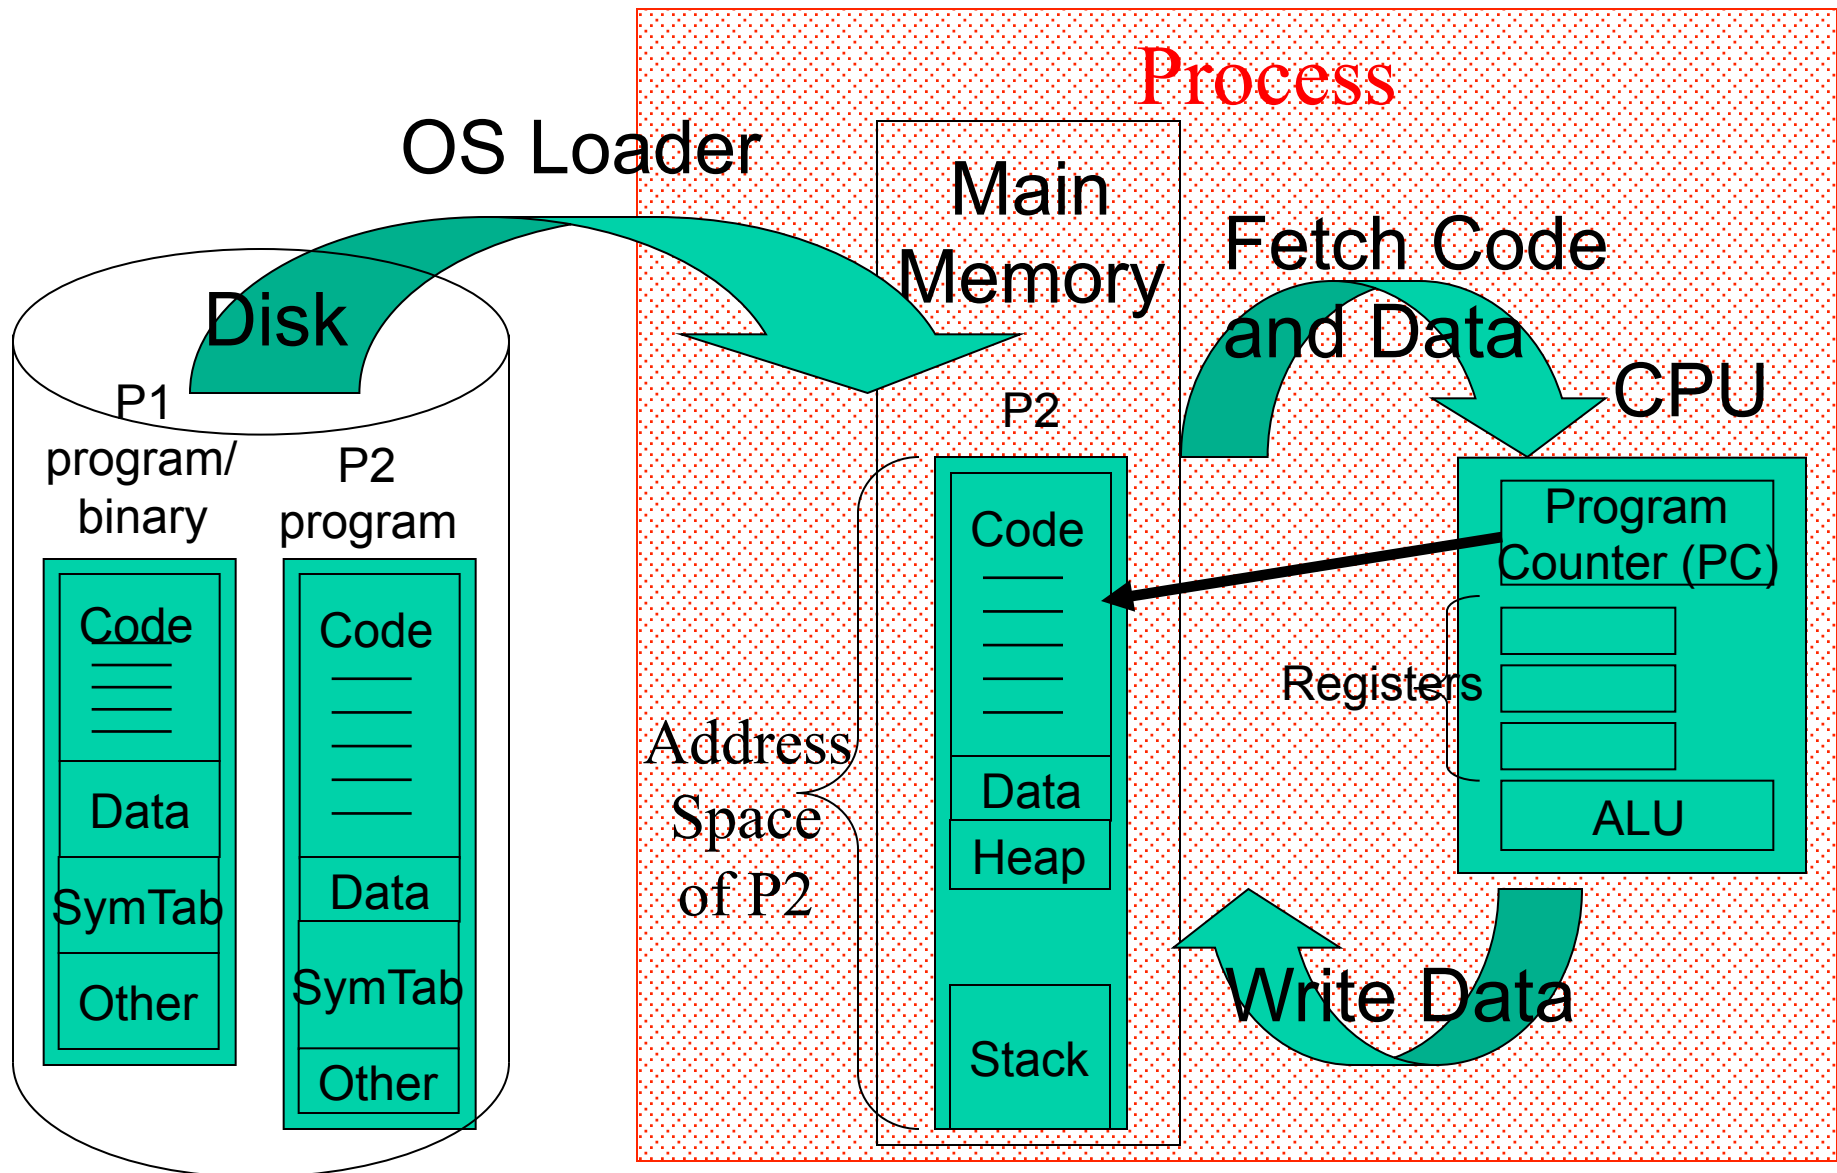


# Chapter 3: What is a Process?

- A software *program* consist of a sequence of code instructions and data stored on disk
  - A program is a *passive* entity
- A *process* is a program *actively* *executing* from main memory within its *own address space*

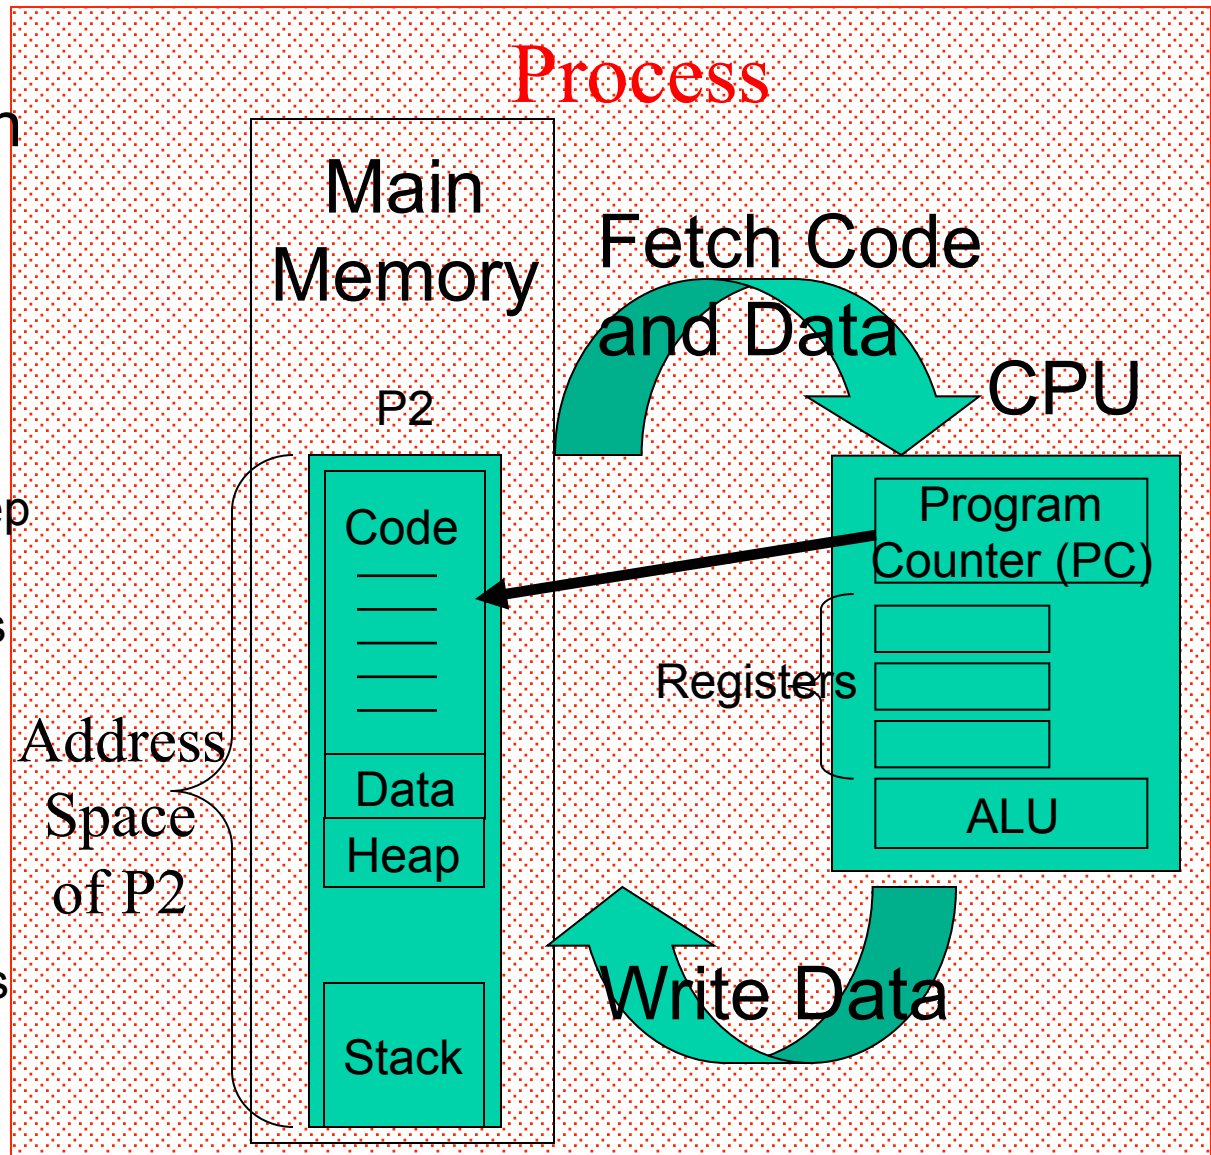


# What Is a Process? (2)



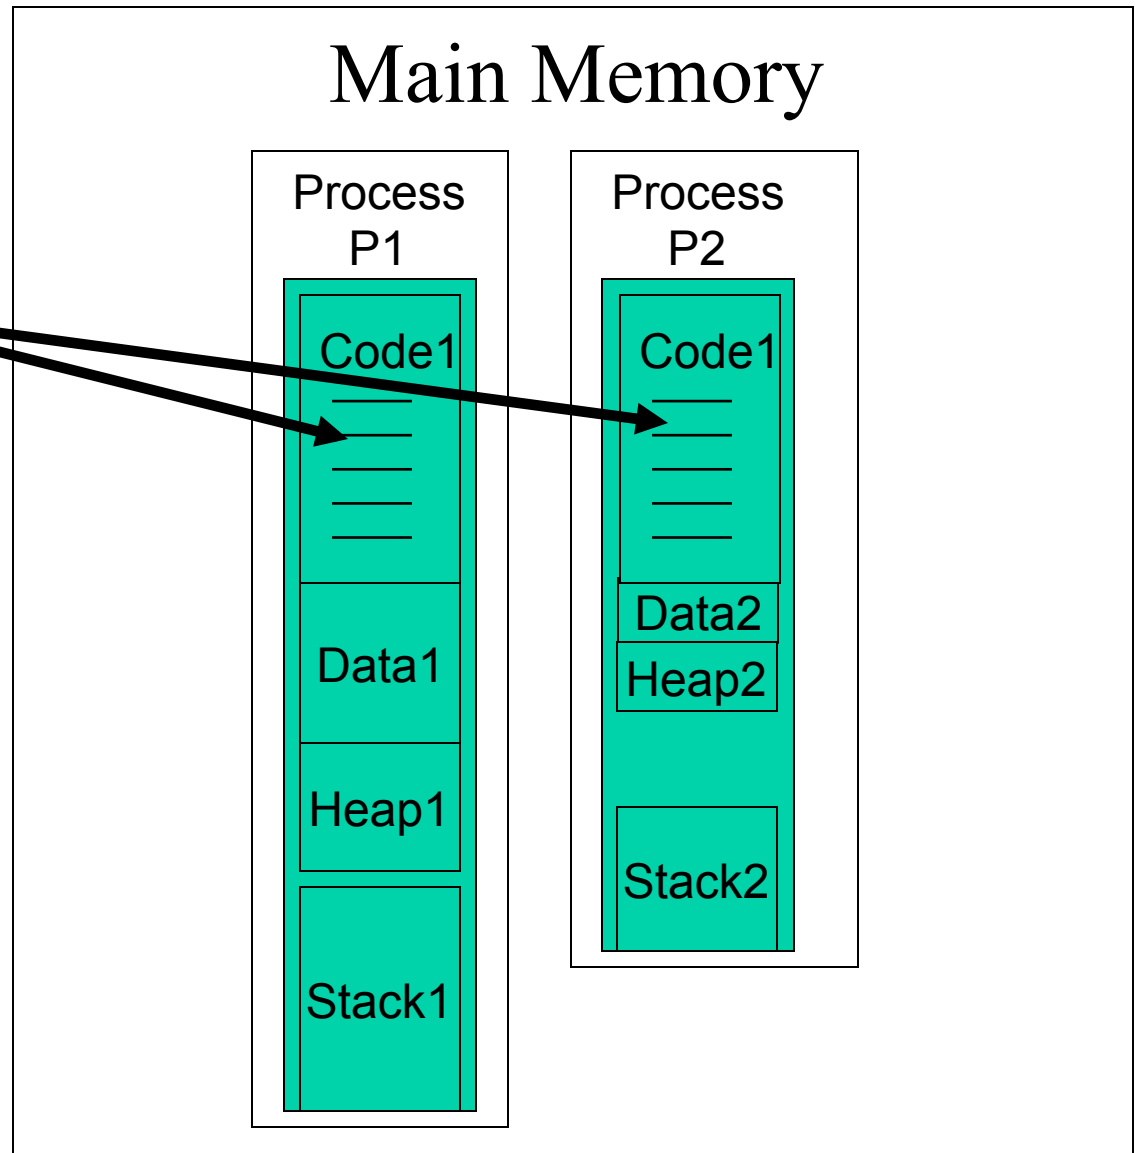
# What is a Process? (3)

- A *process* is a program *actively executing* from main memory
  - has a Program Counter (PC) and execution state associated with it
    - CPU registers keep state
    - OS keeps process state in memory
    - it's alive!
  - Owns its own *address space*
    - a limited set of (virtual) addresses that can be accessed by the executing code



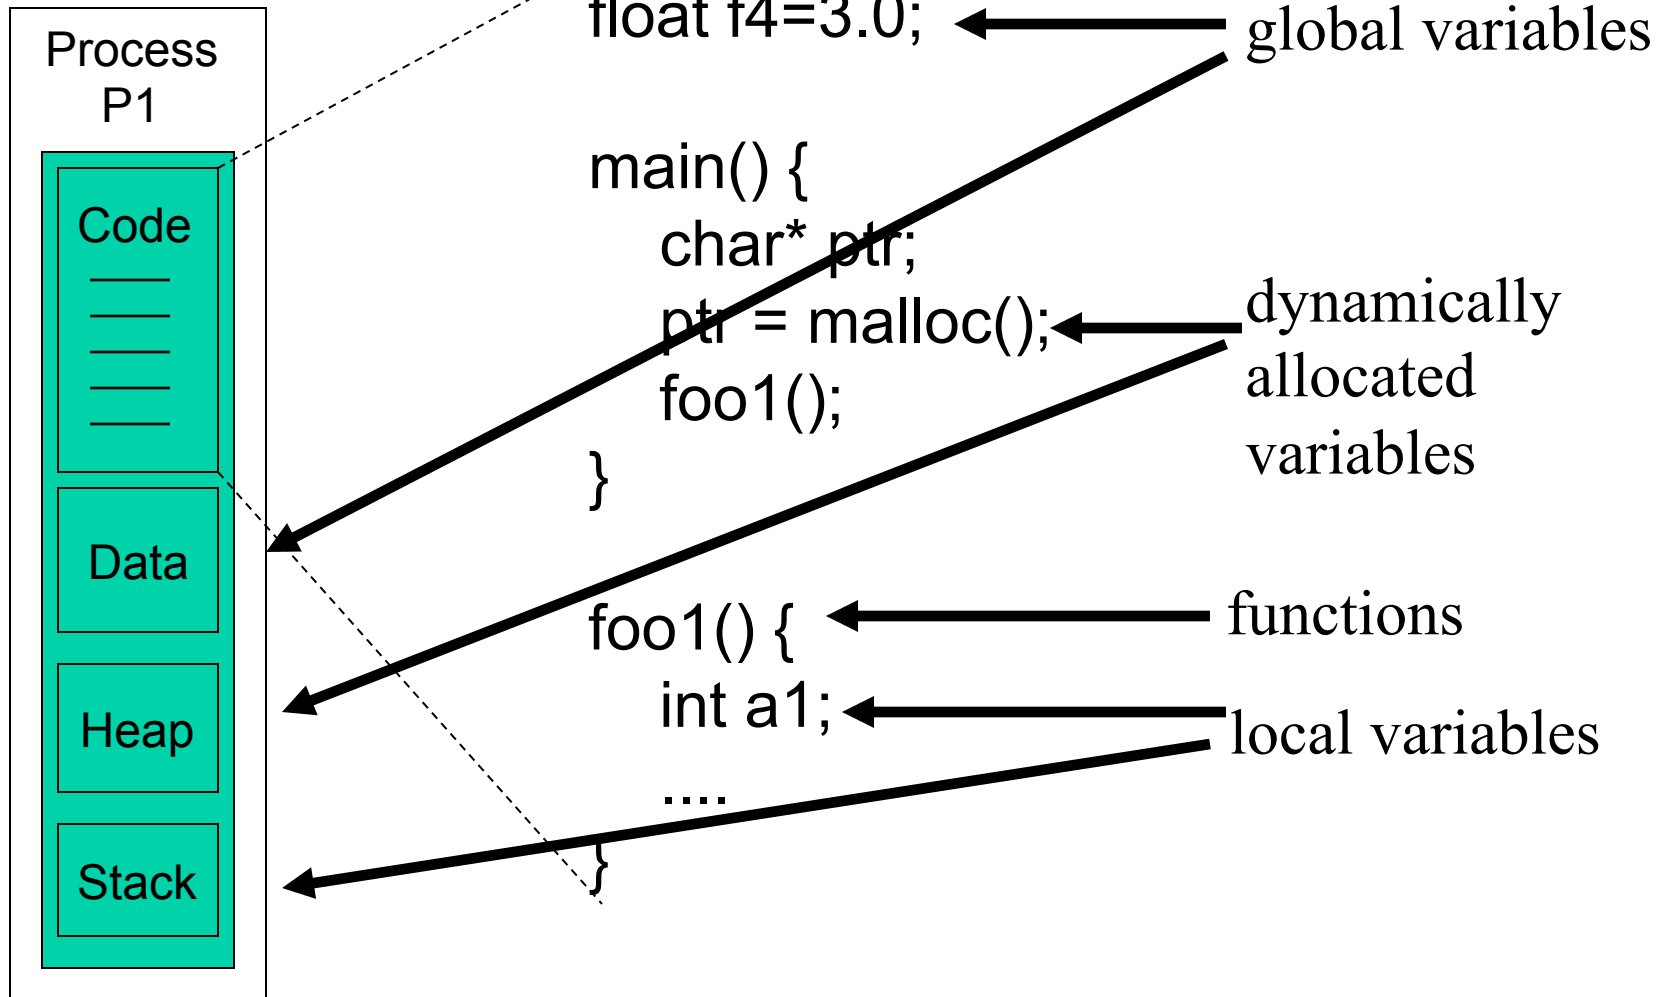
# What is a Process? (4)

- 2 processes may execute the same program code, but they are considered *separate execution sequences*
  - e.g. two shell terminals



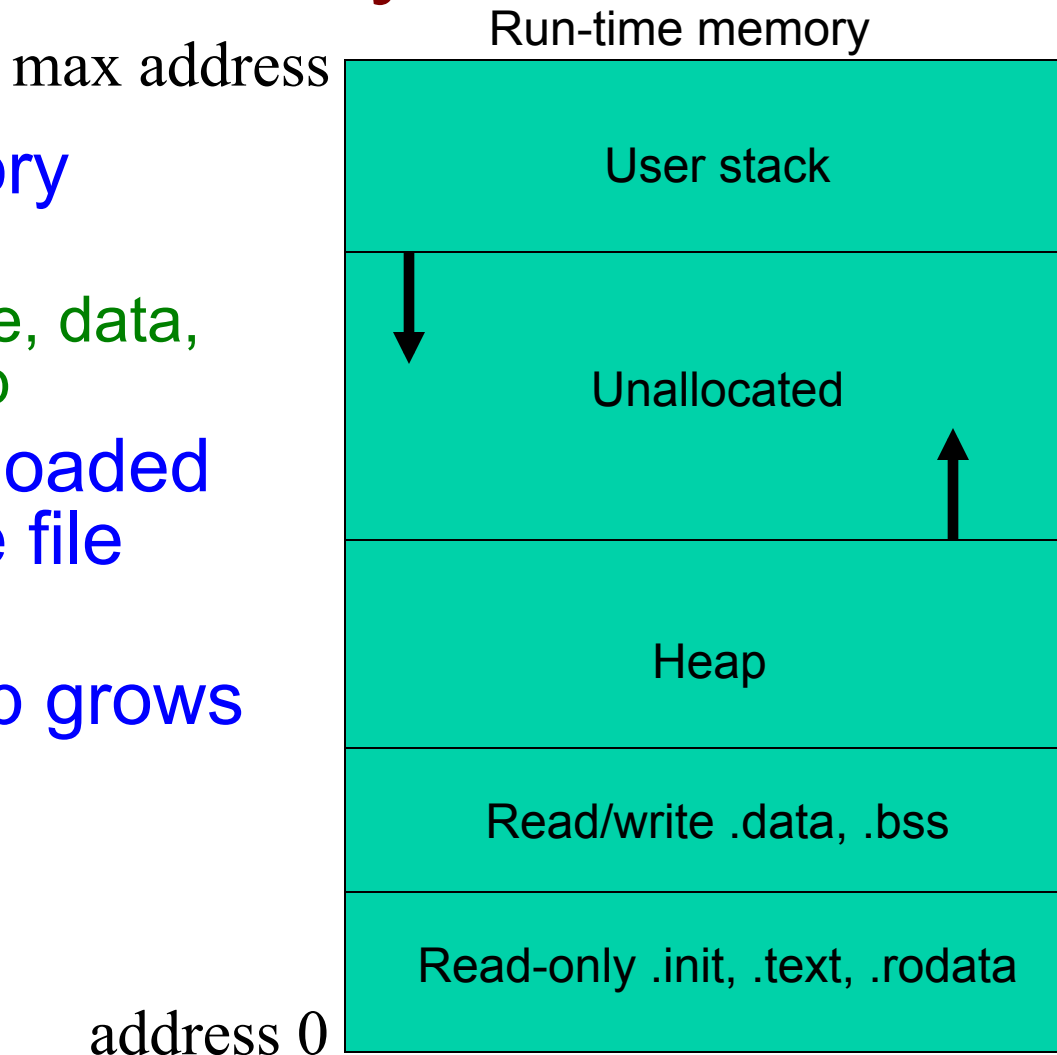
# How is a Process Structured in Memory?

Main Memory




# How is a Process Structured in Memory?

- Run-time memory image
  - Essentially code, data, stack, and heap
- Code and data loaded from executable file
- Stack grows downward, heap grows upward



# A Process Executes in its Own Address Space

Main Memory



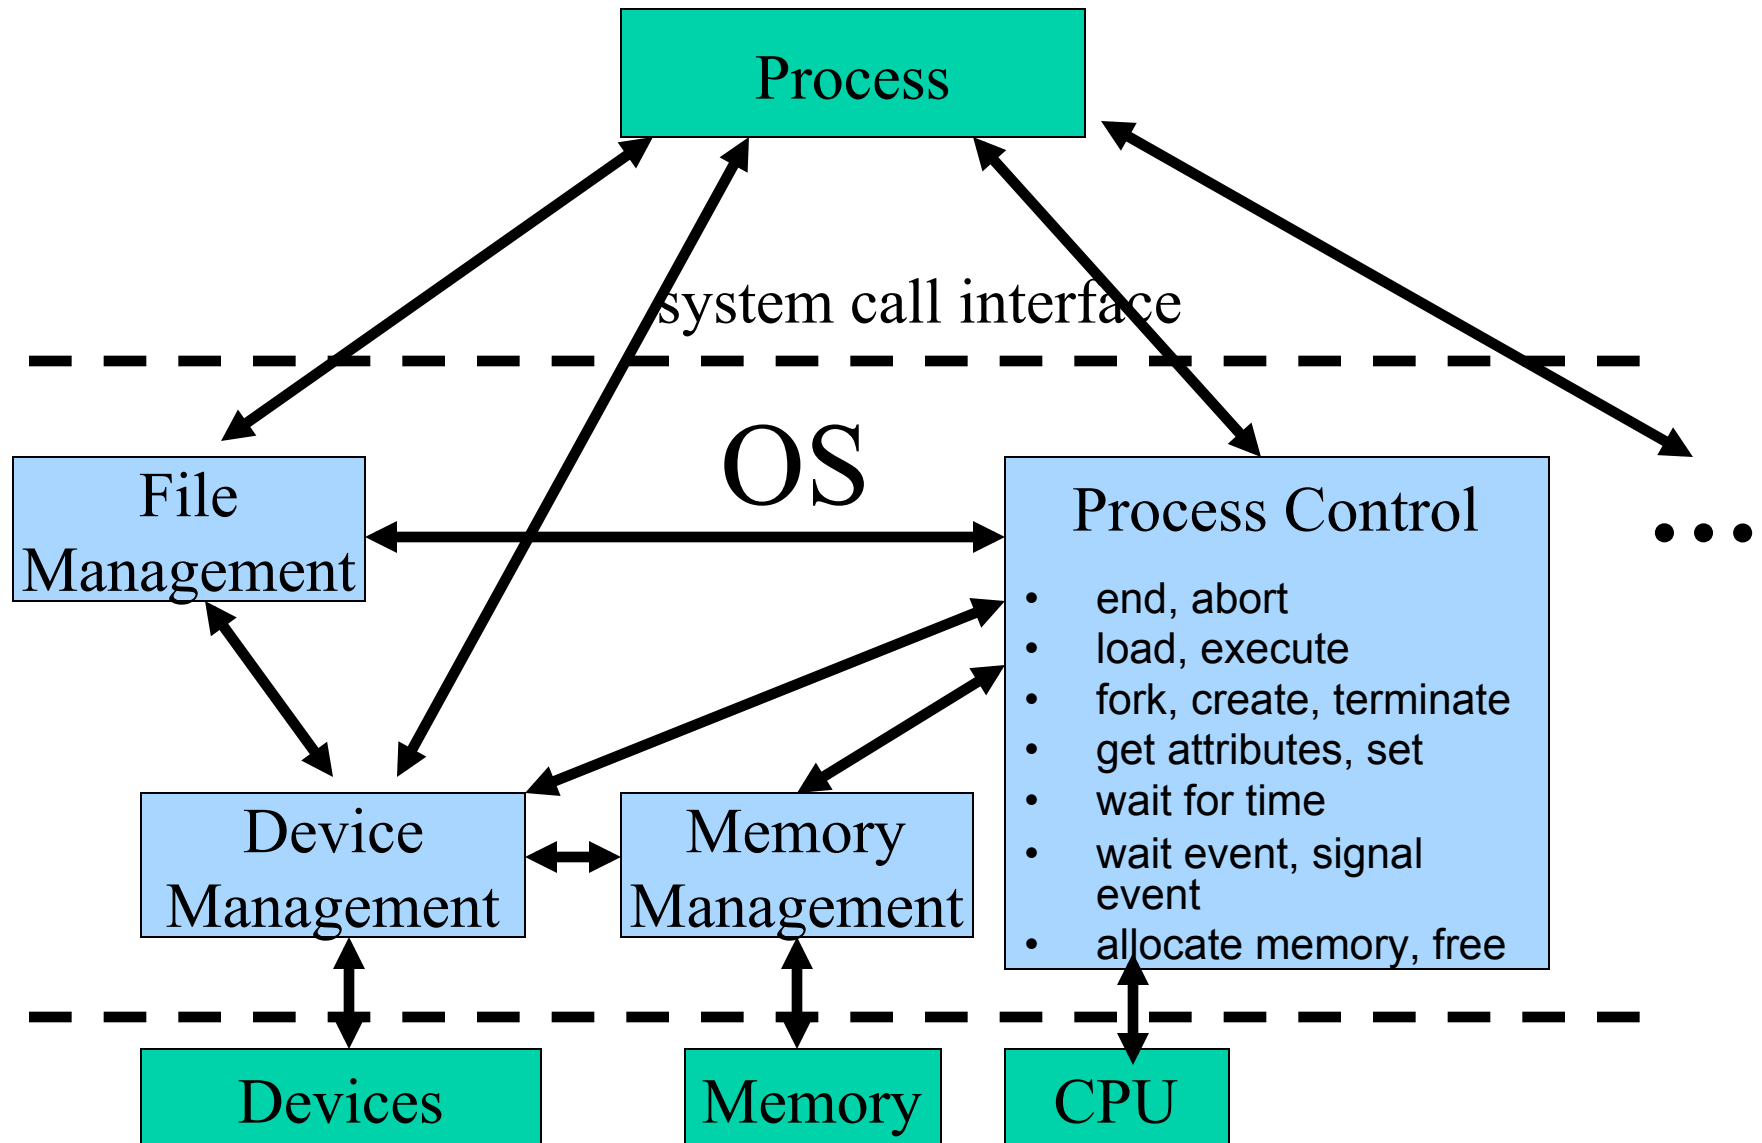
- OS tries to provide the illusion or *abstraction* to the process that it executes
  - in its own subset of RAM, i.e. its own address space
  - on its own subset (time slice) of the CPU

# Applications and Processes

- Application =  $\sum_i$  Processes<sub>i</sub>
  - e.g. a server could be split into multiple processes, each one dedicated to a specific task (UI, computation, communication, etc.)
  - The Application's various processes talk to each other using Inter-Process Communication (IPC). We'll see various forms of IPC later.



# Process Management



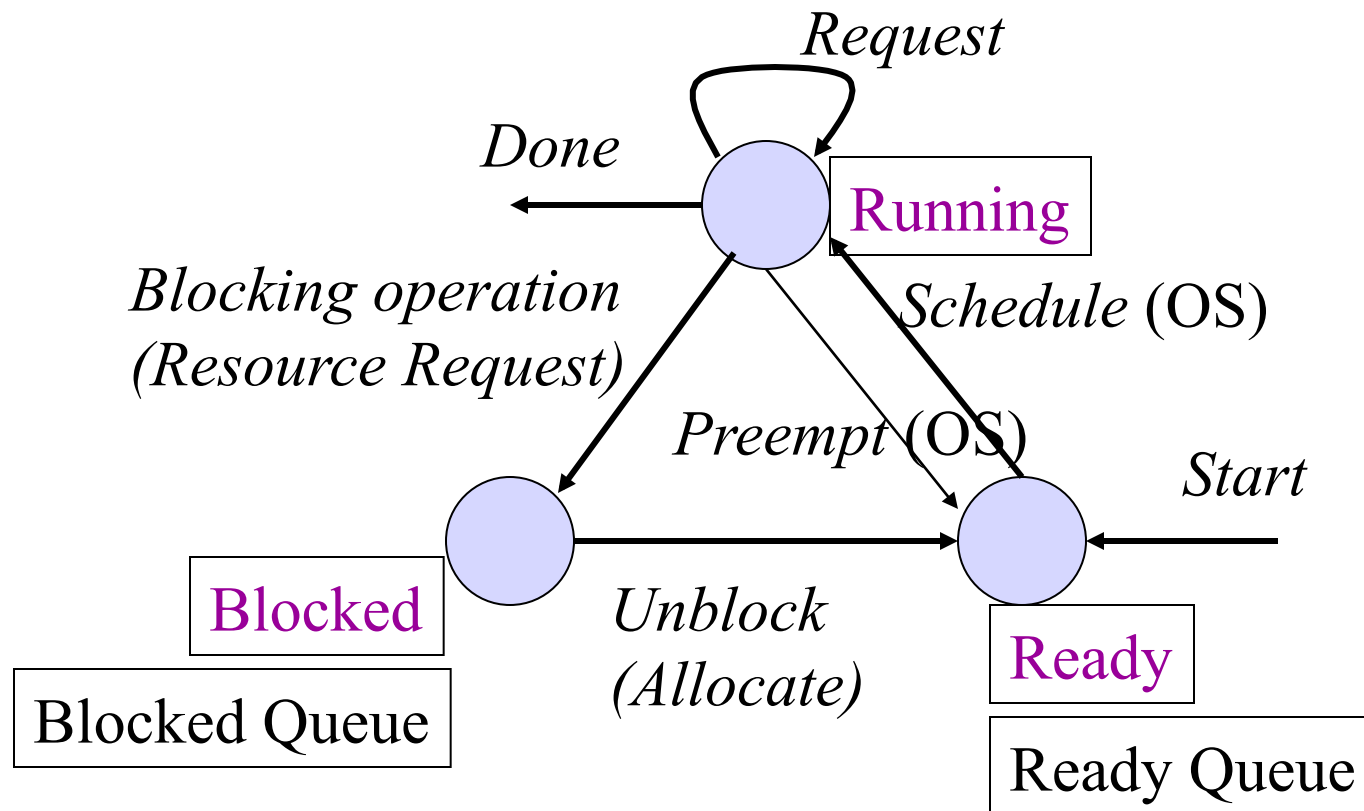
# Process Manager

- Creation/deletion of processes (and threads)
- Synchronization of processes (and threads)
- Managing process state
  - Processor state like PC, stack ptr, etc.
  - Resources like open files, etc.
  - Memory limits to enforce an address space
- Scheduling processes
- Monitoring processes
  - Deadlock, protection

# Process State

- Memory image: Code, data, heap, stack
- Process state, e.g. ready, running, or waiting
- Accounting info, e.g. process ID
- Program counter
- CPU registers
- CPU-scheduling info, e.g. priority
- Memory management info, e.g. base and limit registers, page tables
- I/O status info, e.g. list of open files

# Process State Diagram

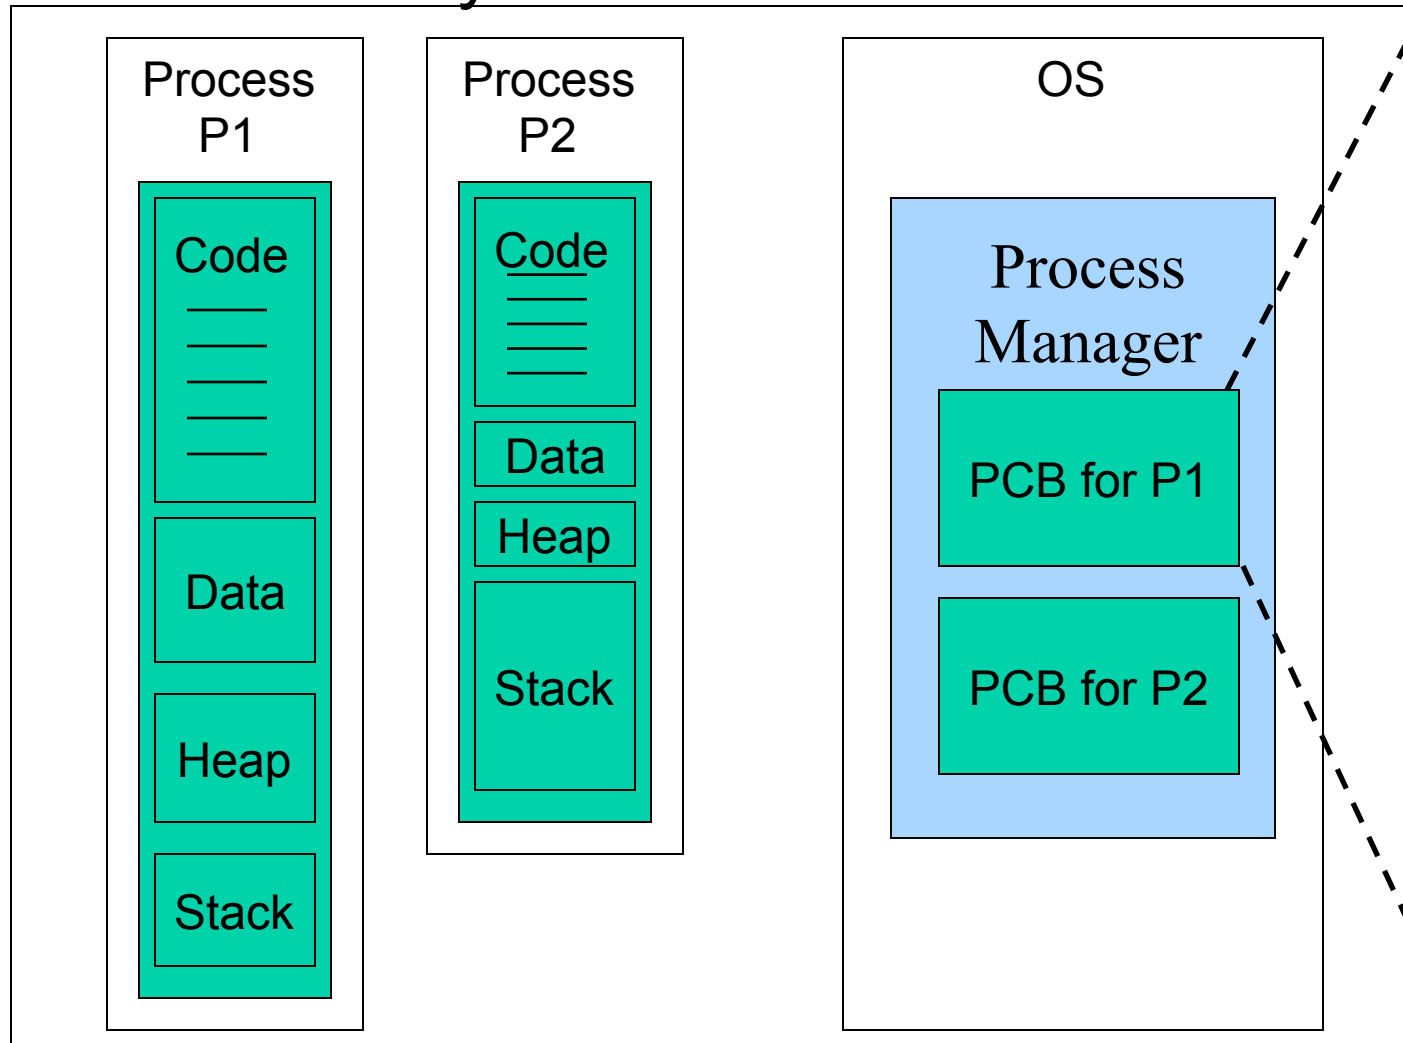


# Process Control Block

- Each process is represented in OS by a process control block (PCB).
- PCB: Complete information of a process
- OS maintains a PCB table containing one entry for every process in the system.
- PCB table is typically of fixed size. This size determines the maximum number of processes an OS can have
  - The actual maximum may be less due to other resource constraints, e.g. memory.

# Process Control Block (PCB)

## Main Memory



- Process state, e.g. ready, running, or waiting
- accounting info, e.g. process ID
- Program Counter
- CPU registers
- CPU-scheduling info, e.g. priority
- Memory management info, e.g. base and limit registers, page tables
- I/O status info, e.g. list of open files

A PCB is also called a Process Descriptor

# Context Switch

- Running state → Ready state
- Running state → Blocked state
- Switching the CPU from currently running process to another process
  - Save the state of the currently running process in its PCB
  - Load the saved state of new process scheduled to run from its PCB
  - Context switch time is pure overhead: 1 – 1000 microseconds.
  - An important goal in OS design is to minimize context switch time.

# Creating Processes

- In Windows, there is a `CreateProcess()` call
  - Pass an argument to *CreateProcess()* indicating which program to start running
  - Invokes a system call to OS that then invokes process manager to:
    - allocate space in memory for the process
    - Set up PCB state for process, assigns PID, etc.
    - Copy code of program name from hard disk to main memory, sets PC to entry point in *main()*
    - Schedule the process for execution
  - As we will see, this combines UNIX's *fork()* and *exec()* system calls and achieves the same effect



# Creating Processes in UNIX

- Use *fork()* command to create/spawn new processes from within a process
  - When a process (called parent process) calls *fork()*, a new process (called child process) is created
  - Child process is an exact copy of the parent process
    - All addresses are appropriately mapped – We'll see this later under memory management
  - The child starts executing at the same point as the parent, namely just after returning from the *fork()* call

# *fork ()*

- The `fork()` call returns an *int* value
  - In the parent process, returned value is child's PID
  - In the child, returned value is 0
  - Since both parent and child execute the same code starting from the same place, i.e. just after the `fork()`, then to differentiate the child's behavior from the parent's, you could add code:

```
PID = fork();  
if (PID==0) { /* child */  
    codeforthechild();  
    exit(0);  
}  
/* parent's code here */
```

# Loading Processes

- The `exec()` system call loads program code into the calling process's memory (same address space!), clears the stack, and begins executing the new code at its main entry point
  - The calling code is erased!
  - Use `fork()` and `exec()` (actually `execve()`) to create a new process executing a new program in a new address space

```
PID = fork();
if (PID==0) { /* child */
    exec("/bin/ls");
    exit(0);
}
/* the parent's code here */
```

# More on Processes

- Copying the entire code of a parent into a child can be expensive on a *fork()*, so
  - at start, child can share parent's code pages. Only create a copy on a write.
- The *wait()* system call is used by a parent process to be informed of when a child has completed, i.e. called *exit()*
  - Once the parent has called *wait()*, the child's PCB and address space can be freed

More about this in recitation

# Process Hierarchy

- OS creates a single process at the start up.
- An existing process can spawn one or more new processes during execution
  - Parent-child relationship
  - A parent process may have some control over its child process(es): suspend/activate execution; wait for termination; etc.
- A tree-structured hierarchy of processes
- Process hierarchy in Unix

*<Figure>*

# Context Switch

- Linux allocates two stacks for each process: a user stack that resides in the user address space and a kernel stack that resides in the kernel
  - Kernel stack is used when the process is executing in the kernel (supervisor mode)
  - Kernel stack is needed for security purposes
  - OS allocates 8 KB for each kernel stack
  - Process's PCB is actually stored at one end of this space and the (kernel) stack starts from the other end
- Linux OS provides *schedule( )*, which is invoked by the timer interrupt to schedule a new process on the CPU
- The *schedule( )* function calls another function *switch\_to( )*, which does the actual context switching

# schedule and switch\_to( ) function

- Here is an outline of schedule and switch\_to functions  
(*actual implementations are in assembly code*)

```
schedule( )
```

```
{  
    disable_interrupts;  
    prev_proc = process id of running process  
    next_proc = process id of the next process to run  
    update ready_queue, system_queue, etc.  
    switch_to (prev_proc, next_proc);  
    running_process = prev_proc;  
    enable_interrupts;  
}
```

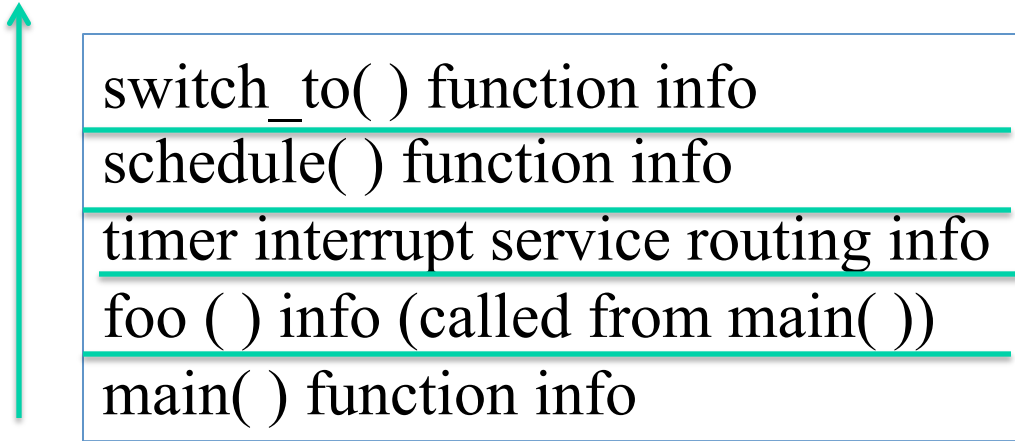
```
switch_to (prev_proc, next_proc)
```

```
{  
    save the state of prev_proc in prev_proc's kernel stack  
    load the state of next_proc from next_proc's kernel stack  
}
```

# Context switch Example

- For simplicity, we will assume a single stack
- Suppose P1 and P2 alternate their execution on CPU
- Lets assume P1 is running and a timer interrupt occurs to preempt this process

P1's stack



A vertical stack of five frames, each separated by a horizontal line. To the left of the stack is a green arrow pointing upwards. The frames contain the following text from top to bottom: 'switch\_to( ) function info', 'schedule( ) function info', 'timer interrupt service routing info', 'foo ( ) info (called from main( ))', and 'main( ) function info'.

switch_to( ) function info
schedule( ) function info
timer interrupt service routing info
foo ( ) info (called from main( ))
main( ) function info

```
switch_to (p1, p2)
{
    save p1's state
    load p2's state
}
```

*P1's instruction ptr*

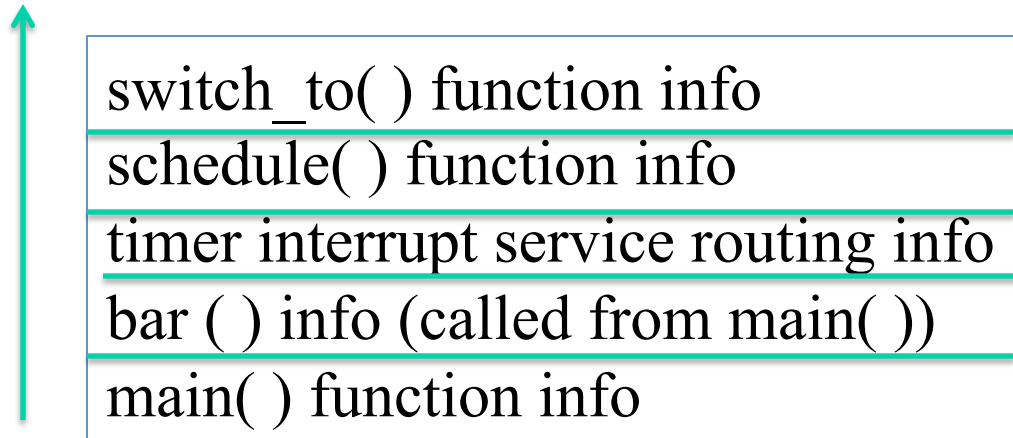


A green arrow points from the text 'P1's instruction ptr' to the closing curly brace of the 'switch\_to' function.




- Since P2's state is loaded, P2 starts running
- Lets assume a timer interrupt occurs to preempt this process

P2's stack



```
switch_to (p2, p1)
{
    save p2's state
    load p1's state
}
```

*P2's instruction ptr*



- Now P1 starts running, since its state is loaded
- P1's instruction pointer is pointing to the end of its `switch_to ( )`, so that function will return
- Next, last two statements of the `schedule` function will execute and that function will return
- Next, the timer interrupt service routine completes and returns
- Finally, the execution of the `foo( )` function resumes from the point where it was executing before P1 was preempted