

# **CSCI 3753**

## **Operating Systems**

**Classic Synchronization Problems**  
**(Producer-Consumer,**  
**Readers-Writers,**  
**Dining Philosophers)**

**Lecture Notes By**  
**Shivakant Mishra**  
**Computer Science, CU-Boulder**  
**Last Update: 02/13/14**

# Classic Synchronization Problems

- Bounded Buffer Producer-Consumer Problem
  - We have already seen this
- Readers-Writers Problem
- Dining Philosophers Problem
- These are not just abstract problems
  - They are representative of several classes of synchronization problems commonly encountered in the real world when trying to synchronize access to shared resources among multiple processes or threads

# Readers writers problem

- A database is accessed by two types of processes: reader processes and writer processes.
- Readers only read information from the database.
- Writers modify the database.
- Constraints
  - Writers must have exclusive access to the database.
  - Multiple readers can access the database concurrently.

# Readers-Writers Problem: First Attempt

Semaphore mutex = 1;

Reader( )

```
{  
    wait(mutex);  
    read database  
    signal(mutex);  
}
```

Writer( )

```
{  
    wait(mutex);  
    write database  
    signal(mutex);  
}
```

Exclusive access to the writer processes is provided  
BUT: No concurrency among reader processes

# Readers-Writers Problem: Second Attempt

```
int rc = 0; /* Number of readers in the database */
```

```
Semaphore db = 1; /* controls access to database for writers */
```

```
Semaphore mutex = 1; /* controls access to variable rc */
```

```
Reader ()
```

```
    wait(mutex);
```

```
    if (rc == 0) wait(db);
```

```
    rc++;
```

```
    signal(mutex);
```

```
        read database
```

```
    wait(mutex);
```

```
    rc--;
```

```
    if (rc == 0) signal(db);
```

```
    signal(mutex);
```

```
Writer( )
```

```
{
```

```
    wait(db);
```

```
        write database
```

```
    signal(db);
```

```
}
```

# Readers-Writers Problem: Second Attempt

- Semaphore mutex is used for mutual exclusion to update rc
- Semaphore db is used for exclusive database access for writer processes
- Multiple reader processes can access the database concurrently if there is no writer process.

Problem: What happens to a writer if readers keep coming to read the database?

- Writer starvation
- Readers have priority over writers.


- Write a solution that gives preference to writer processes
- Write a solution that is fair to both readers and writers

# Dining Philosophers Problem

- The most famous synchronization problem.
- Represents a situation that can occur in large community of processes that share a large pool of resources.
- Five philosophers sit around a round dining table. A plate of spaghetti is placed in front of each philosopher, and a fork is placed between any two adjacent plates.
- A philosopher needs two forks to eat spaghetti.
- All philosophers alternate between two activities: thinking and eating.




# Dining Philosophers Problem



Write a synchronization program that allows all five philosophers to run their lives

- Deadlock free
- Starvation free



# Dining Philosophers: First Attempt

```
philosopher(int i) {  
    while(TRUE) {  
        // Think  
        wait(fork[i]);  
        wait(fork[(i+1) mod 5]);  
        // EAT  
        signal(fork[(i+1) mod 5]);  
        signal(fork[i]);  
    }  
}  
semaphore fork[5] = (1,1,1,1,1);
```

## Problem

- On a fateful day, all philosophers decide to eat at the same time.
- All philosophers pick up their right fork.
- All philosophers now wait forever for their left fork to become available, *and die of starvation*
  - Deadlock.

# Dining Philosophers: Second Attempt

```
philosopher(int i) {  
    while(TRUE) {  
        // Think  
        wait(mutex);  
        wait(fork[i]);  
        wait(fork[(i+1) mod 5]);  
        signal(mutex);  
        // Eat  
        signal(fork[(i+1) mod 5]);  
        signal(fork[i]);  
    }  
}  
  
semaphore fork[5] = (1,1,1,1,1);  
semaphore mutex = 1;
```

- This solution doesn't suffer from deadlock
- Problem
  - May not allow two non-adjacent philosophers to eat at the same time

# Dining Philosophers: Third Attempt

After picking up a fork, if a philosopher finds that the other fork is not available, she keeps down the fork, waits for some time, and tries again.

- Does this solution work?
- Starvation
- A deadlock-free solution is not necessarily starvation-free

# Dining Philosophers: Some possible solutions

- Allow at most 4 philosophers at the same table when there are 5 resources
- Odd philosophers pick first left then right, while even philosophers pick first right then left
- Allow a philosopher to pick up forks *only if both are free*. This requires protection of critical sections to test if both forks are free before grabbing them.
  - we'll see this solution next using monitors
  - Also, there is a construct called an AND semaphore



# Higher Level Synchronization Primitives

- Semaphores can result in deadlock due to programming errors
  - Forgot a wait() or signal(), or mis-ordered them, or duplicated them
- Relatively simple problems, such as the dining philosophers problem, can be very difficult to solve using low level constructs like semaphores
- Higher level synchronization primitives
  - AND synchronization
  - Events
  - Critical Conditional Regions
  - Condition Variables: *We will study this*
  - Monitors: *We will study this*
  - many others...

# Pthreads: Condition Variables

- Provide support for synchronization between two or more threads
- Used with a `pthread_mutex_t` variable
- Six operations
  - `pthread_cond_init ( )`: Initialization
  - `pthread_cond_wait ( )`
    - Always blocks
    - Releases mutex before blocking
    - Re-acquires mutex before returning
    - Re-acquiring the mutex can block the thread for a little longer

- `pthread_cond_signal ( )`
  - wakes up *at least one* thread blocked on the condition variable
- `pthread_cond_broadcast ( )`
  - wakes up *all* of the threads blocked on the condition variable
- `pthread_cond_timedwait ( )`
  - Identical to `pthread_cond_wait()`, except it has a timeout
  - This timeout is an absolute time of day
  - If timeout has expired, it returns ETIMEDOUT
- `pthread_cond_destroy ( )`
  - Deallocation

# Monitors

- Abstract data type (similar to C++ classes)
  - Monitors are found in high-level programming languages like Java and C#
- A monitor is a collection of procedures, variables, and data structures
- Processes can access these variables only by calling procedures in the monitor
- Each function in the monitor can only access variables declared locally within the monitor and its parameters
- At most one process may be active at any time in a monitor

- monitor *monitor\_name* {  
    // shared local variables

```
function f1(...) {
```

```
...
```

```
}
```

```
...
```

```
function fN(...) {
```

```
...
```

```
}
```

```
init_code(...) {
```

```
...
```

```
}
```

```
}
```

# Monitors and Condition Variables

- While the above definition of a monitor achieves mutual exclusion (hiding wait() and signal() from user), it loses the ability that semaphores had to enforce order
  - i.e. wait() and signal() are used to provide mutual exclusion, but the unique ability for one process to signal another blocked process using signal() is lost
- In general, there may be times when one process wishes to signal another process based on a condition, much like semaphores.
  - Thus, augment monitors with *condition variables*.

# Monitors and Condition Variables

- Declare a condition variable with pseudo-code  
    condition x, y;
- A condition variable x in a monitor allows three main operations on itself
  - x.wait()
    - blocks the calling process
    - can have multiple processes suspended on a condition variable, typically released in FIFO order, but textbook describes another variation specifying a priority p, i.e. call x.wait(p)
  - x.signal()
    - resumes exactly 1 suspended process. If none, then *no effect*.
  - x.queue()
    - Returns true if there is at least one process blocked on x

- Note that `x.signal()` is unlike the semaphore's signaling operation `V()`, which preserves state in terms of the value of the semaphore.
  - Example: if a process `Y` calls `x.signal()` on a condition variable `x` before process `Z` calls `x.wait()`, then `Z` will wait. The condition variable doesn't remember `Y`'s signal.
  - Comparison: if a process `Y` calls `signal(mutex)` on a binary semaphore `mutex` (initialized to 0) before process `Z` calls `wait(mutex)`, then `Z` will not wait, because the semaphore remembers `Y`'s `signal()` because its value = 1, not 0.



# Monitors and Condition Variables

- Within a monitor, if a process P1 calls `x.signal()`, then normally that would wake another process P2 blocked on `x.wait()`. But we must avoid having two processes at the same time in the monitor, so need “wake-up” semantics on a `x.signal()`:
  - Hoare semantics, also called signal-and-wait
    - The signaling process P1 either waits for the woken up process P2 to leave the monitor before resuming, or waits on another CV
  - Mesa semantics, also called signal-and-continue
    - The signaled process P2 waits until the signaling process P1 leaves the monitor or waits on another condition

# Dining Philosophers: Monitor-based Solution

- Key insight: pick up 2 forks only if both are free
  - Avoids deadlock
  - A philosopher moves to his/her eating state only if both neighbors are not in their eating states
    - Need to define a state for each philosopher
  - If one of my neighbors is eating, and I'm hungry, ask them to signal() me when they're done
    - States of each philosopher: thinking, hungry, eating
    - Need condition variables to signal() waiting hungry philosopher(s)
  - Also, need to Pickup() and Putdown() forks

# Dining Philosophers: Monitor-based Solution

```
monitor DiningPhilosophers
```

```
{
```

```
    enum {Thinking, Hungry, Eating} state[5];
```

```
    condition self[5];
```

```
    void test(int i) {
```

```
        //Called by philosopher i or neighbors of i
```

```
        //Check if both neighbors of i are not eating
```

```
        //If so, set state[i] to Eating, and signal philosopher i
```

```
    }
```

```
    void pickup(int i) {
```

```
        //Called by philosopher i
```

```
        //Set state[i] to Hungry and call test(i)
```

```
        //If at least one neighbor is eating, block on self[i];
```

```
    }
```

*... cond. to the next slide*

*... cond. from the previous slide*

```
void putdown(int i) {  
    //Called by philosopher i after eating  
    //change state[i] to Thinking and signal neighbors in  
    //case they are waiting to eat  
}
```

```
init( ) {  
    for (int i = 0; i < 5; i++)  
        state[i] = Thinking;  
}
```

```
philosopher (int i)  
{  
    while (1) {  
        //Think  
        DiningPhilosophers.pickup(i);  
        // pick up forks and eat  
        DiningPhilosophers.putdown(i);  
    }  
}
```

```
void test(int i) {  
    if ((state[(i+1)%5] != Eating) &&  
        (state[(i-1)%5] != Eating) &&  
        (state[i] == Hungry))  
    {  
        state[i] = eating;  
        self[i].signal();  
    }  
}
```

```
void pickup(int i) {  
    state[i] = hungry;  
    test(i);  
    if(state[i]!=Eating)  
        self[i].wait;  
}
```

```
void putdown(int i) {  
    state[i] = thinking;  
    test((i+1)%5);  
    test((i-1)%5);  
}
```

# Starvation

- Note that starvation is still possible in the DP monitor solution
  - Suppose P1 arrives first, and start eating, then P2 arrives and sets its state to hungry and blocks
  - Next P3 arrives and starts eating
  - P1 ends, but P2 can't start eating because P3 is eating
  - Now P1 starts eating again before P3 finishes eating
  - P3 ends, but P2 still can't eat
  - P1 and P3 can alternate this way and P2 will never get to eat →starvation