

CSCI 3753

Operating Systems

Memory Management

Paging and Segmentation

Chapters 8 and 9

Lecture Notes By
Shivakant Mishra
Computer Science, CU-Boulder
Last Update: 03/18/14

Paging

- One of the problems with fragmentation is finding a sufficiently large contiguous piece of unallocated memory to fit a process into
 - Heavyweight solution is to compact memory
- Another solution to external fragmentation is to divide the logical address space into fixed-size **pages**
 - Main memory is divided into similarly-sized frames. Each page is mapped to a page frame.
 - Frames allocated to process don't have to be contiguous in memory. Each process is now scattered throughout memory
 - Solves external fragmentation problem
 - Need a **page table** to keep track of where each logical page of a process is located in main memory, i.e. to keep track of the mapping of each logical page to a physical memory frame

Paging

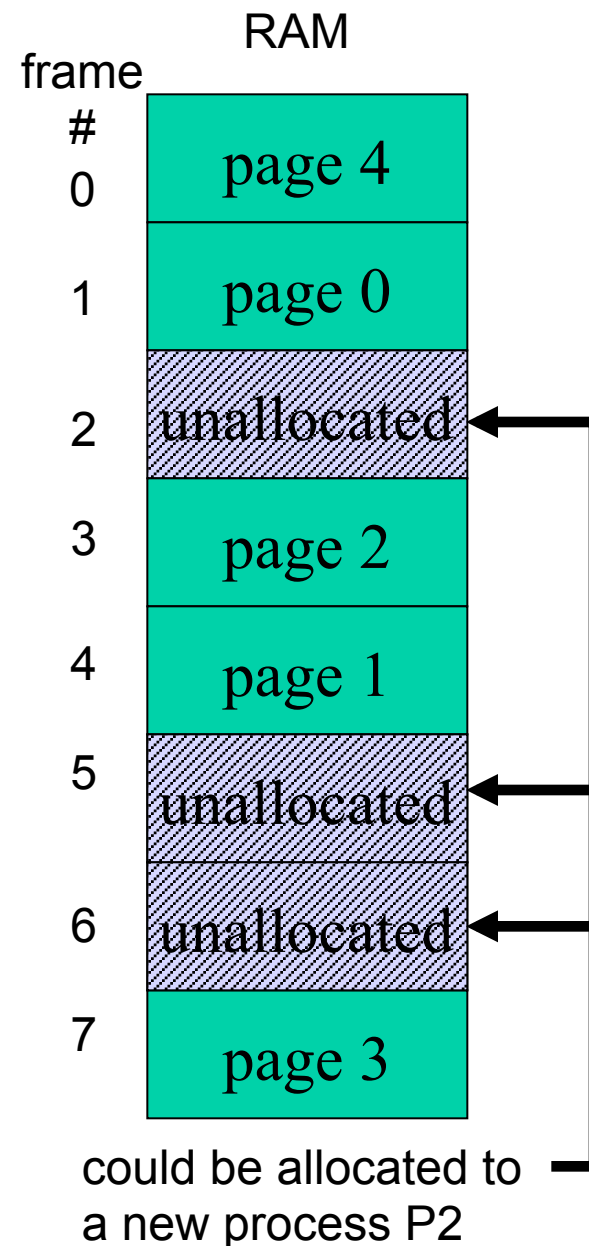
- A page table for each process is maintained by the OS
- Given a logical address, MMU will determine to which logical page it belongs, and then will consult the page table to find the physical frame in RAM to access

Logical Address
Space

page 0
page 1
page 2
page 3
page 4

Page Table

Logical page	Physical frame
0	1
1	4
2	3
3	7
4	0

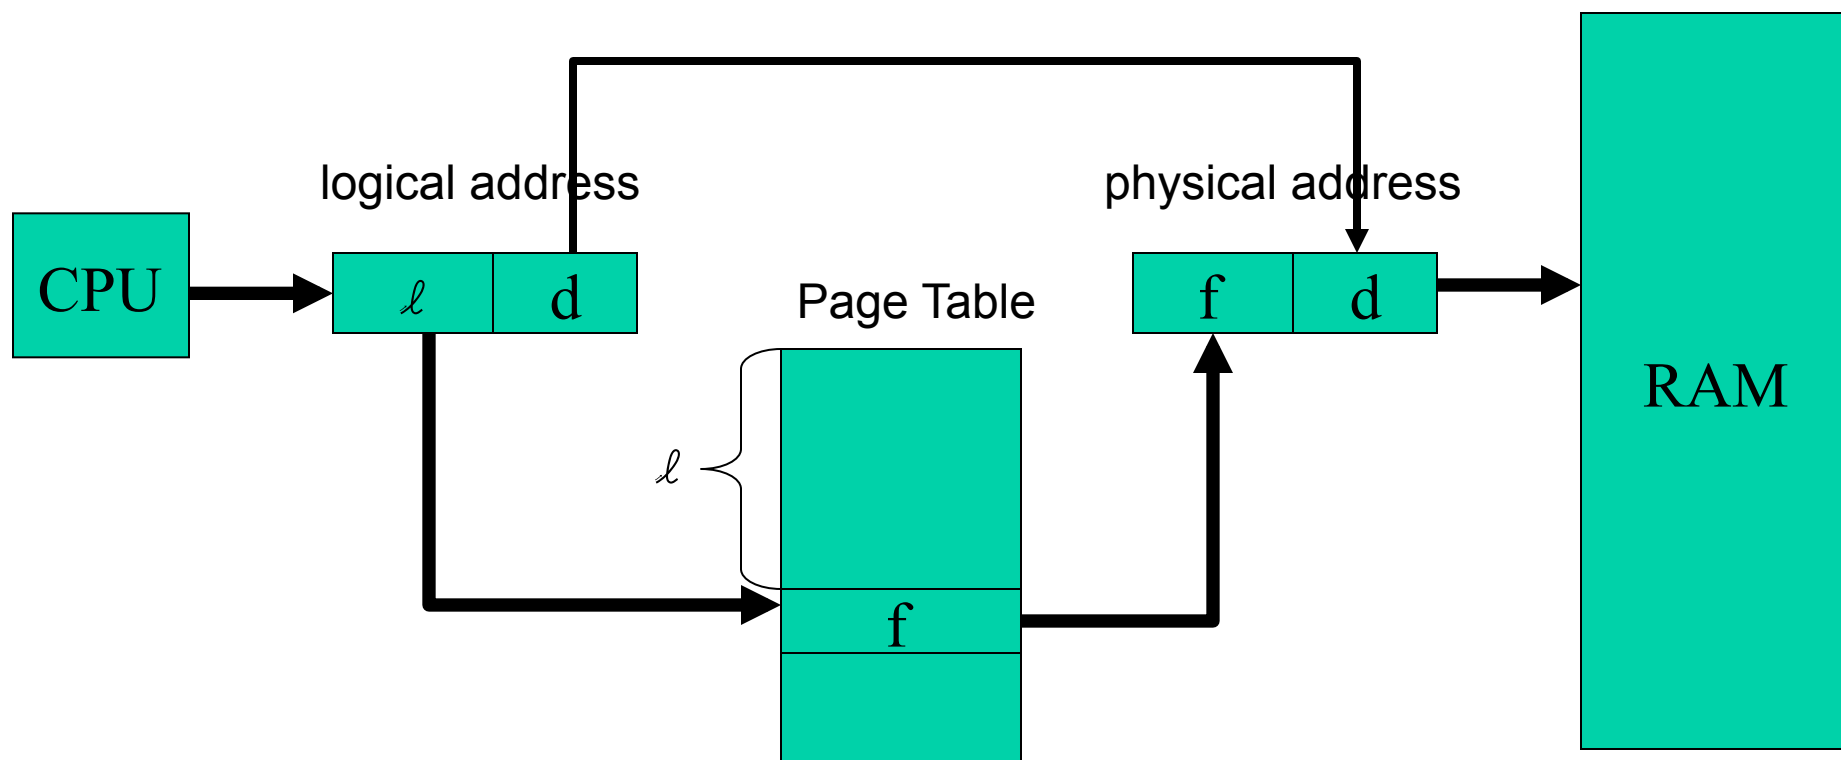


Paging

- Typical page size is 4-8 KB
 - Linux system call: *getpagesize()*; command line: *getconf PAGESIZE*
 - Example: if a page table allows 32-bit entries, and page size is 4 KB, then can address 2^{32} bytes = 16 TB of memory
 - Example: a 4 GB 32-bit address space with 4 KB/page (2^{12}) implies that there can be $2^{32}/2^{12} = 1$ million entries in a process's page table. Your page table would need to be ≥ 20 bits/entry.
- No external fragmentation, but internal fragmentation
 - Example: if my process is $(2^{12} + 1)$ B, and each page size is 4 KB, then I have to allocate two pages = 8 KB, so that 4095 B of 2nd page is wasted due to fragmentation internal to a page
- OS also has to maintain a frame table that keeps track of what frames are free

Paging

- Conceptually, every logical address can be divided into two parts:
 - most significant bits = logical page # ℓ , used to *index* into page table to retrieve the corresponding physical frame f
 - least significant bits = page offset d



Implementing Page Tables

- Address translation must be very fast
- Option #1: Use dedicated bank of hardware registers or memory to store the page table
 - Fast per-instruction translation
 - Slow per context switch - entire page table has to be reloaded
 - Limited by cost (expensive hardware) to being too small - some page tables can be large, e.g. 1 million entries – too expensive

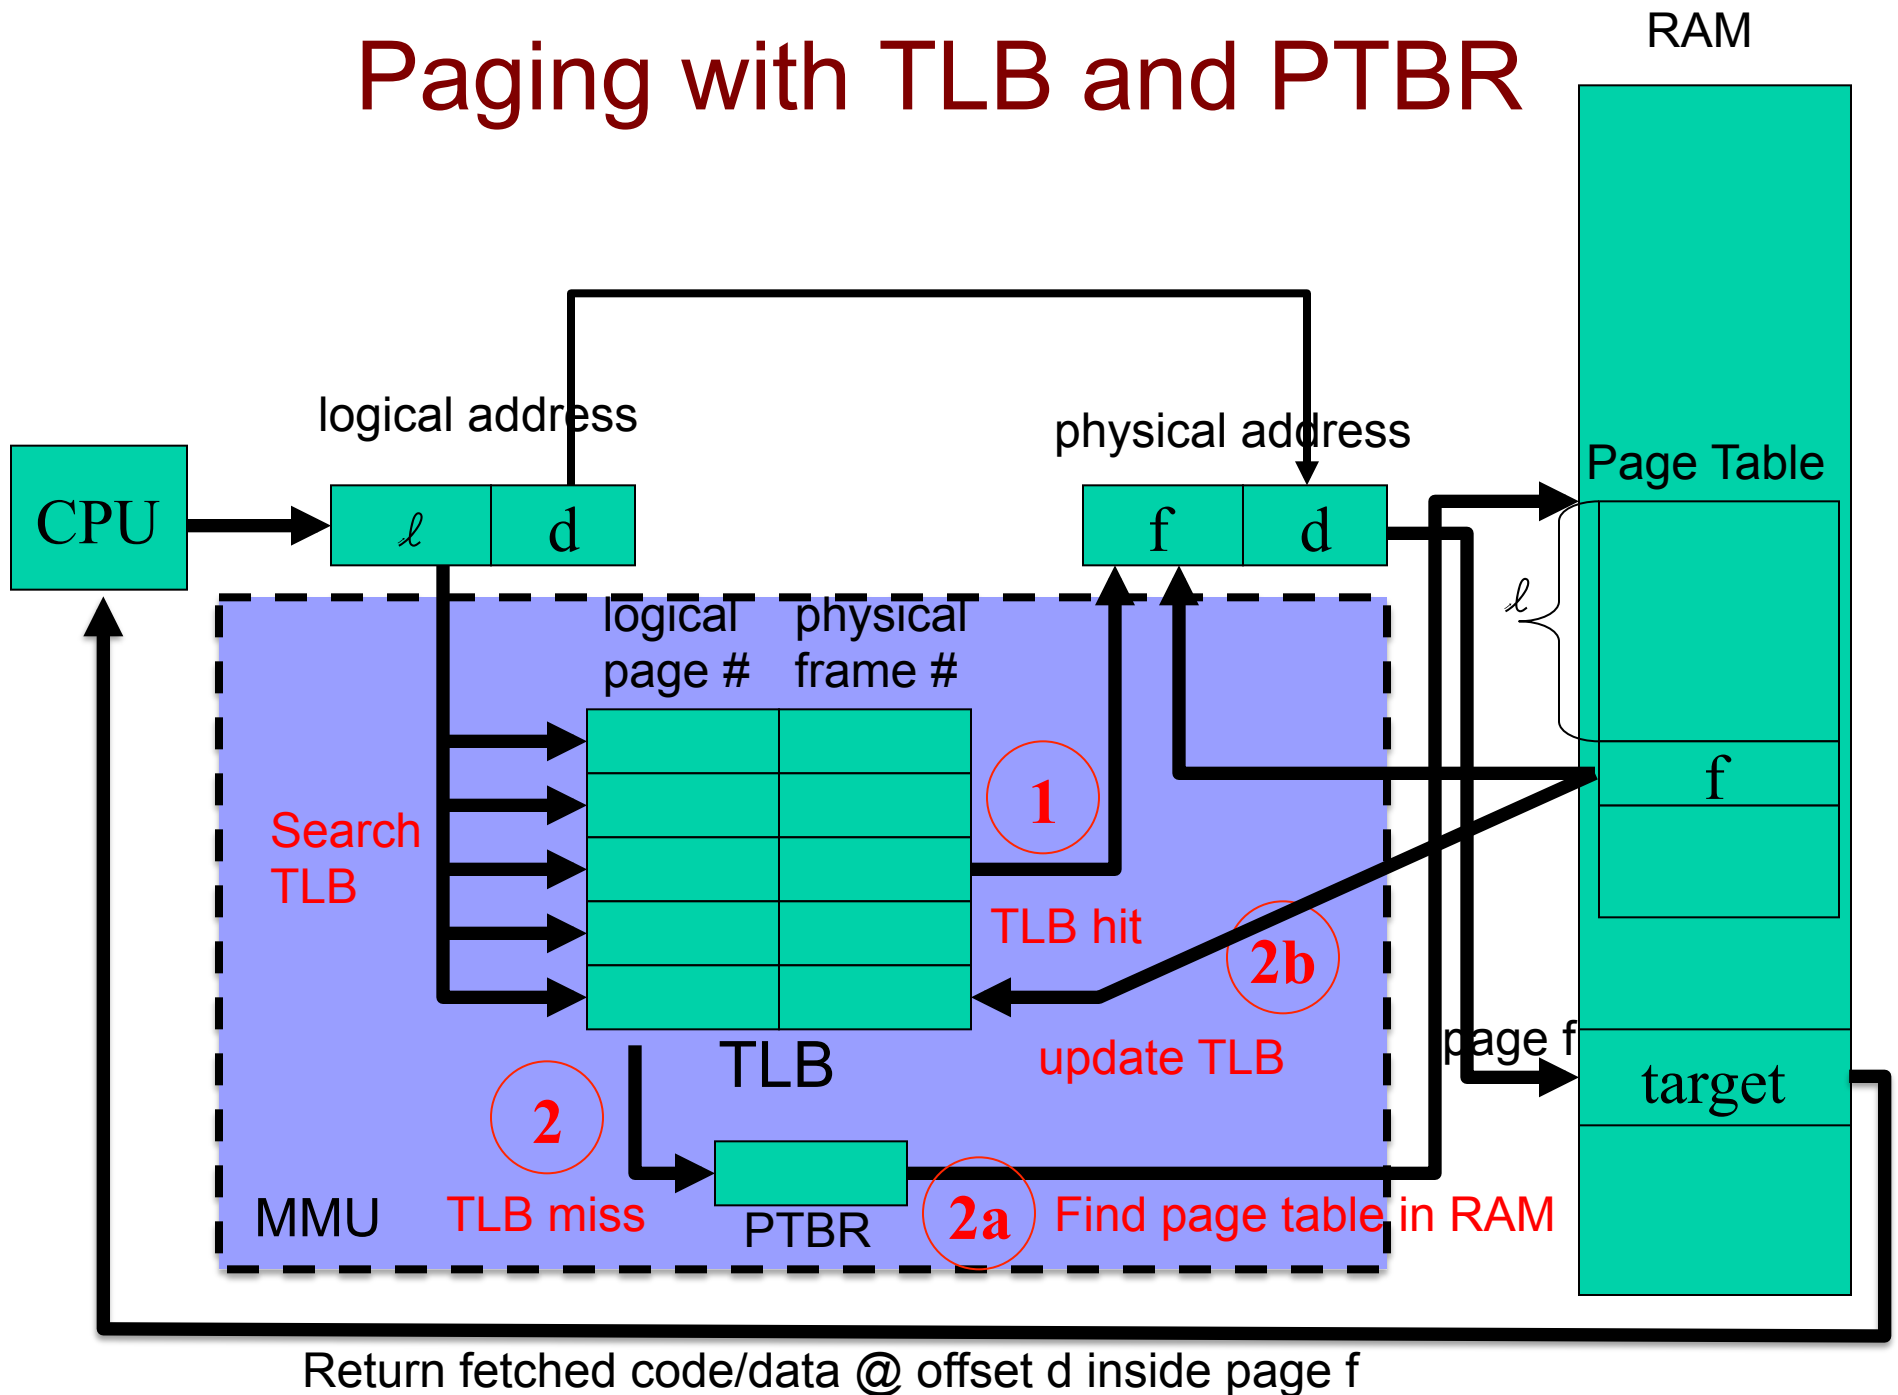
Implementing Page Tables

- Option #2: Store the page table in main memory and just keep a pointer to the page table in a special CPU register called the Page Table Base Register (PTBR)
 - Can accommodate fairly large page tables
 - Fast context switch - only PTBR needs to be reloaded
 - Slow per-instruction translation, because each instruction fetch requires two steps:
 1. finding the page table in memory and indexing to the appropriate spot to retrieve the physical frame # f
 2. retrieving the instruction from physical memory frame f

Paging and Caching

- Option #3:
 - Cache a subset of page table mappings/entries in a small set of CPU buffers called *Translation-Look-aside Buffers* (TLBs)
 - TLB as implemented in hardware does a fast parallel match of the input page to all stored values in the cache - about 10% overhead in speed
 - <key, value> pair
 - Several TLB caching policies
 - Cache the most popular or frequently referenced pages in TLB
 - Cache the most recently used pages
 - Goal is to maximize TLB hits and minimize TLB misses

Paging with TLB and PTBR



Paging and Caching

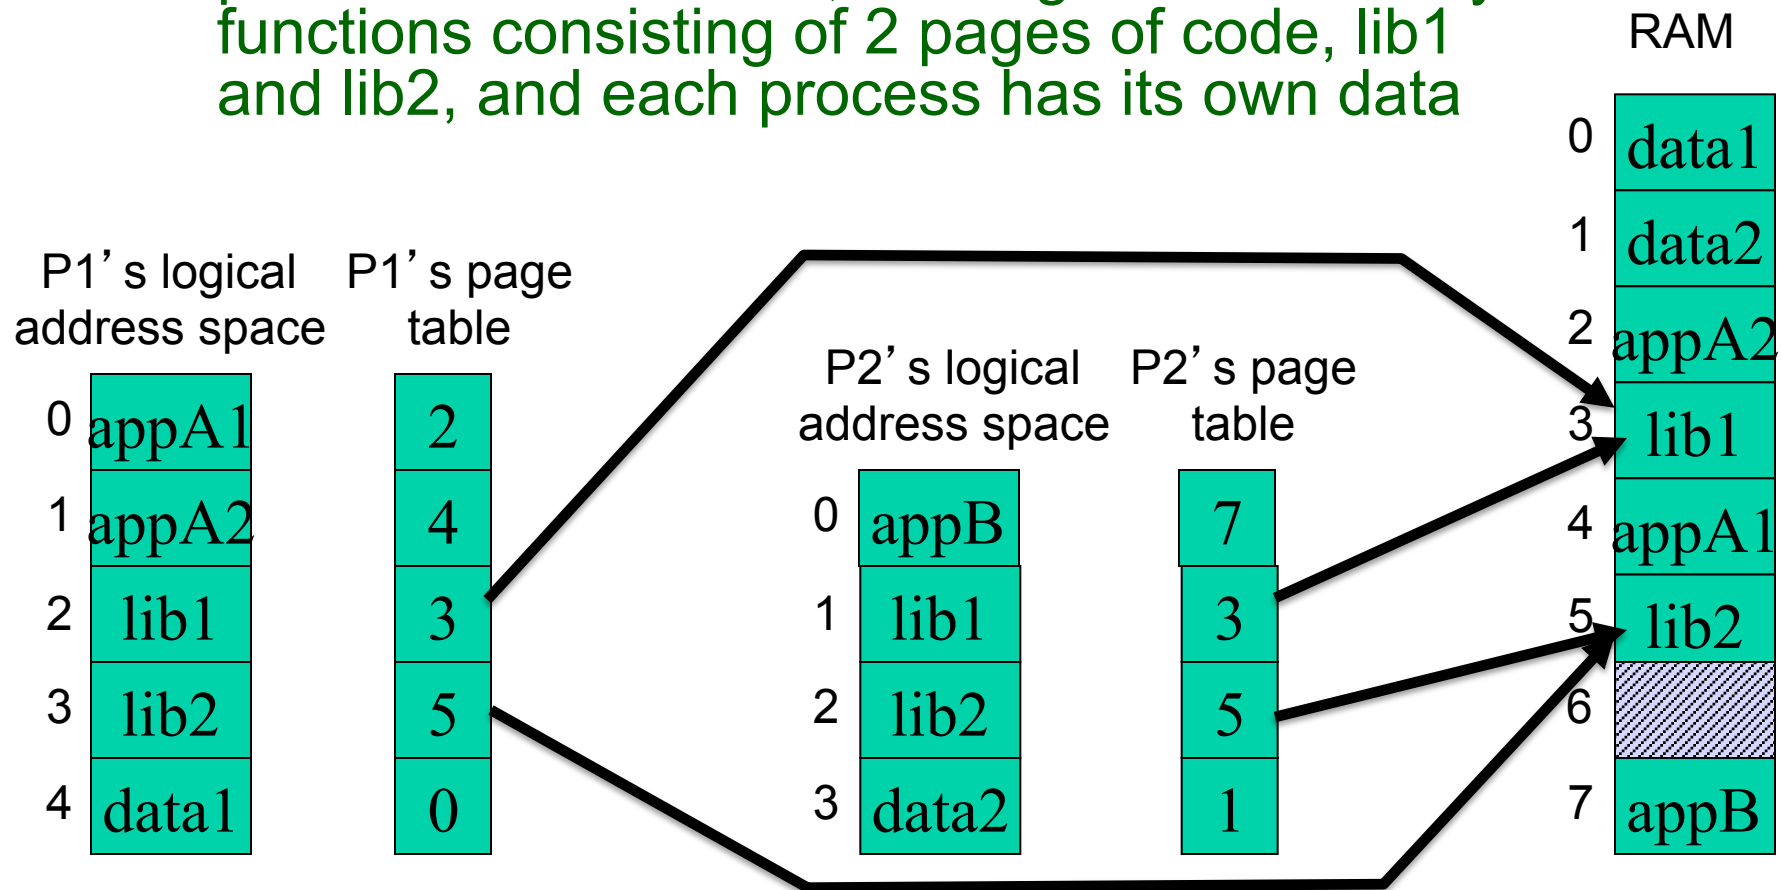
- On a context switch, since different processes have different page tables, the on-chip TLB entries would typically have to be entirely invalidated/completely flushed (x86 behavior)
 - An alternative is to include process IDs in TLB, at the additional cost of hardware and an additional comparison per lookup. Only TLB entries with a process ID matching the current task are considered valid. (See DEC RISC Alpha CPU)
 - In Intel Pentium Pro, the page global enable (PGE) flag in the register CR4 and the global (G) flag of a page-directory or page-table entry can be used to prevent frequently used pages from being automatically invalidated in the TLBs on a task switch
 - ARM allows flushing of individual entries from the TLB indexed by virtual address

Shared Pages

- It is possible and in some cases desirable to share code & data pages between processes by simply having entries in different process' page tables point to the same physical page/frame(s)
 - **Sharing code:**
 - Fork()'ing a child process causes the child to have a copy of the entire address space of the parent, including code. Rather than duplicating all such code pages, can simply map the child's page table to point to the same set of code pages as the parent.
 - May want to share dynamically linked libraries (image, C, ...) between multiple processes, so map each process' page table to point to the same dll pages in memory
 - Shared code should be thread-safe and reentrant
 - **Sharing data:**
 - Two or more processes may want to share memory between them, so pointing multiple page tables to the same data pages is a way to implement shared memory.
 - Shared data should be protected by synchronization

Shared Pages

- To achieve sharing of code and data, page tables can point to the *same* memory frames
 - Consider 2 apps A and B, running in 2 processes P1 and P2, sharing some C library functions consisting of 2 pages of code, lib1 and lib2, and each process has its own data



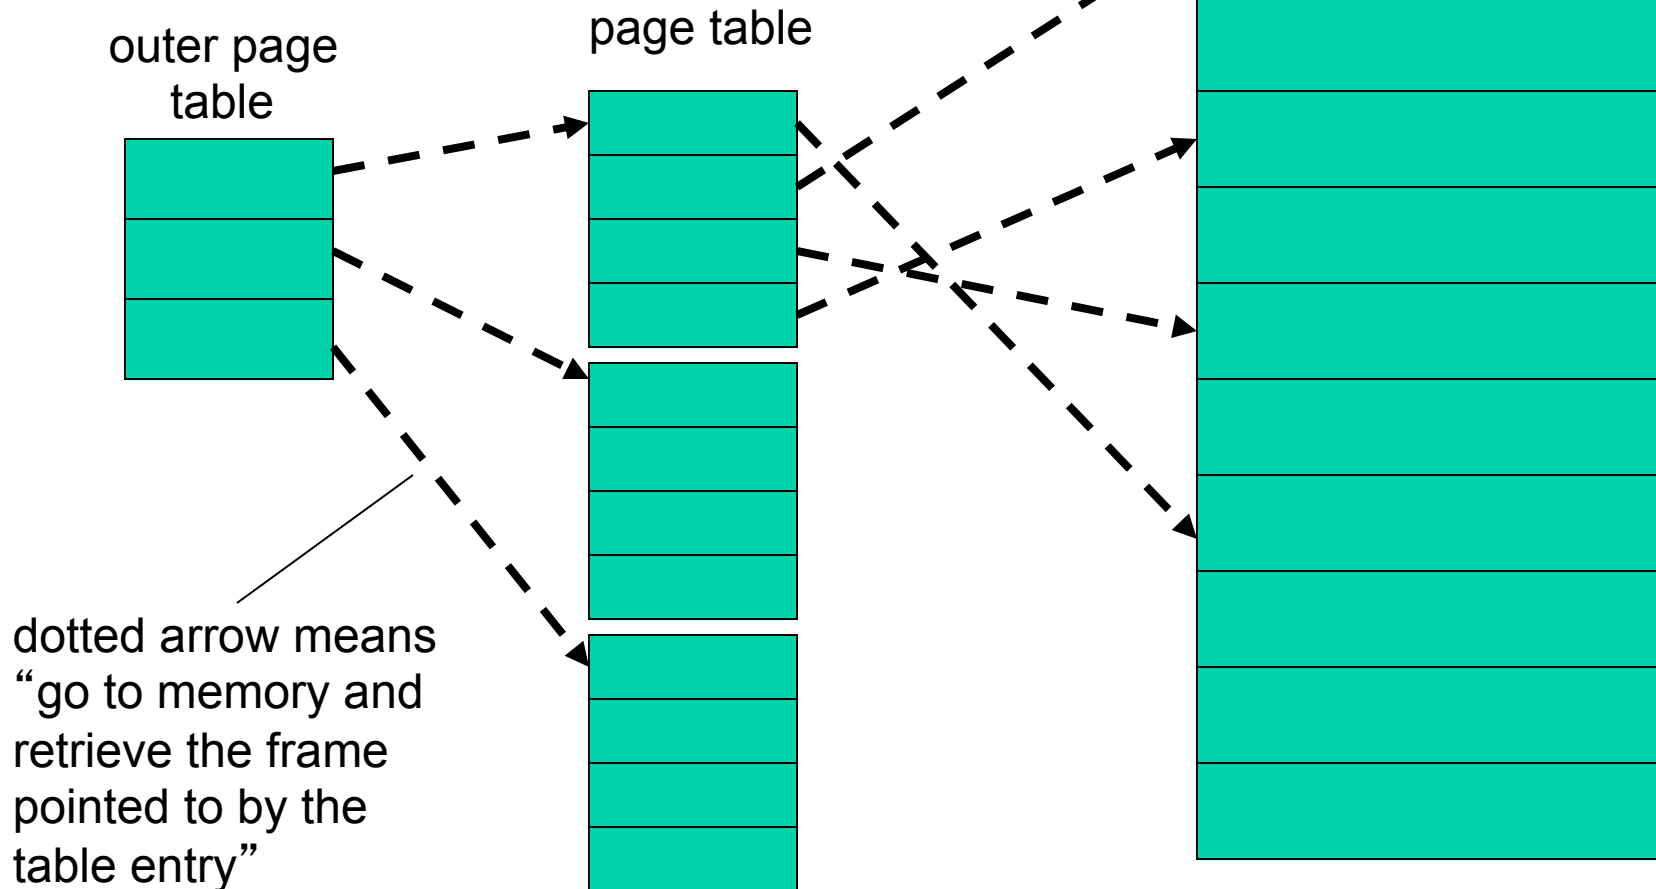
Hierarchical Paging

- Problem with page tables: they can get very large
 - It's hard to find contiguous allocation to store each page table, each of which can be as large as 1 MB.
 - Solution: page the page table → *hierarchical paging*
 - Subdividing a process into pages helped to fit a process into memory by allowing it to be scattered in non-contiguous pieces of memory, thereby solving the external fragmentation problem
 - So reapply that principle here to fit the page table into memory, allowing the page table to be scattered non-contiguously in memory
 - This is an example of 2-level paging. In general, we can apply N-level paging.

Hierarchical Paging

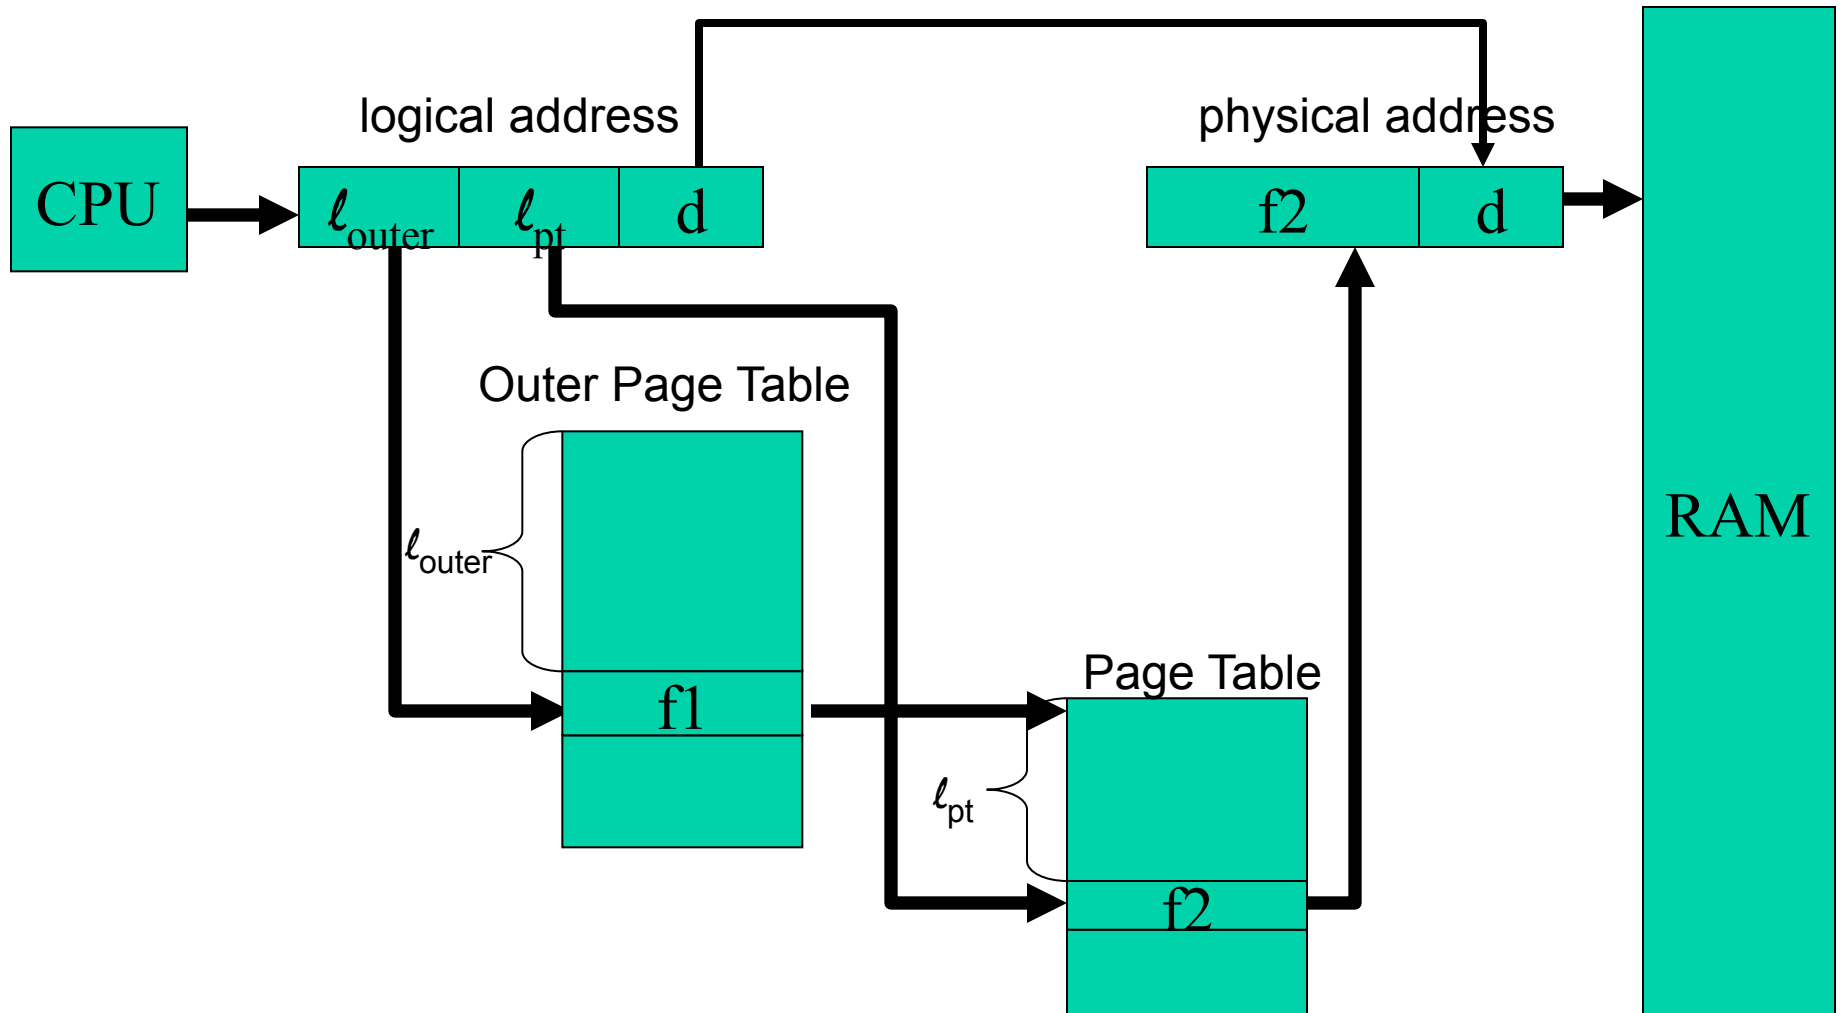
RAM

- Hierarchical (2-level) paging:
 - Outer page table tells where each part of the page table is stored, i.e. indexes into the page table



Hierarchical Paging

- Hierarchical (2-level) paging divides the logical address into 3 parts:



Hierarchical Paging

- Hierarchical page tables do not solve the problem of the large page table size
 - if several processes have a page table that contains a million entries, then a not-insignificant fraction of RAM becomes devoted to storing/managing page tables
 - The page tables may be *sparse* - many of the entries in the page table are just empty placeholders for logical pages that are not in memory, and may never be in memory
 - e.g. stack and heap may never grow large enough to use all of allocated memory

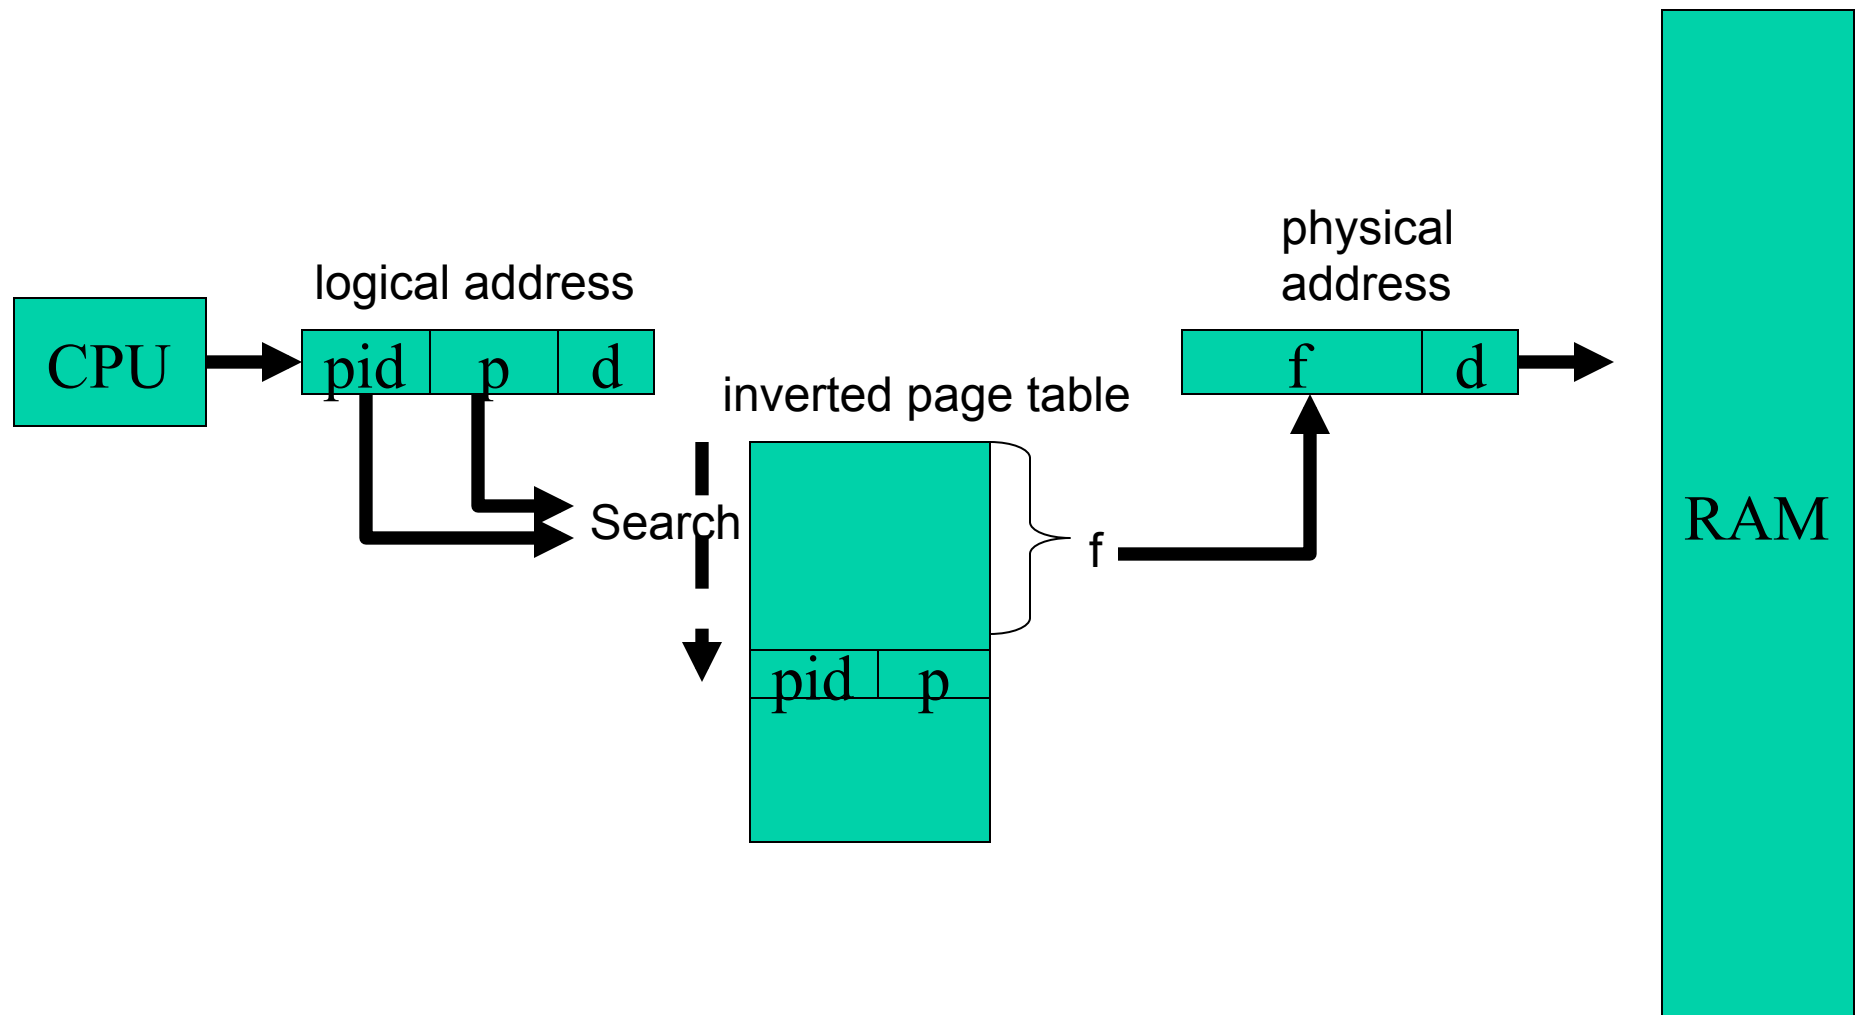
Inverted Page Table

- Solution: Use inverted page table (IPT)
 - Have only one page table for all of memory, rather than one for each process
 - This saves on memory
 - Each entry in the inverted page table lists which process owns a physical frame f , and what logical page # p is stored in that physical frame
 - Each entry in the IPT lists a process id pid and a logical page p , namely $IPT[f] = \langle pid, p \rangle$
 - Thus the index into an inverted page table is the physical frame # f
 - Compare to a page table, whose index is the logical page # p

Inverted Page Table

- Still want to use an inverted page table to map a logical page to a physical frame
 - Since the IPT stores an array of $\langle \text{pid}, p \rangle$ pairs, then given a pid and logical page p , you have to *search* through the IPT to find an entry that matches
 - this can be slow – use hash table
 - It's hard to implement shared memory pages with IPTs, since only one process owns each physical frame, whereas in shared memory multiple processes can use the same code in a physical frame

Inverted Page Table



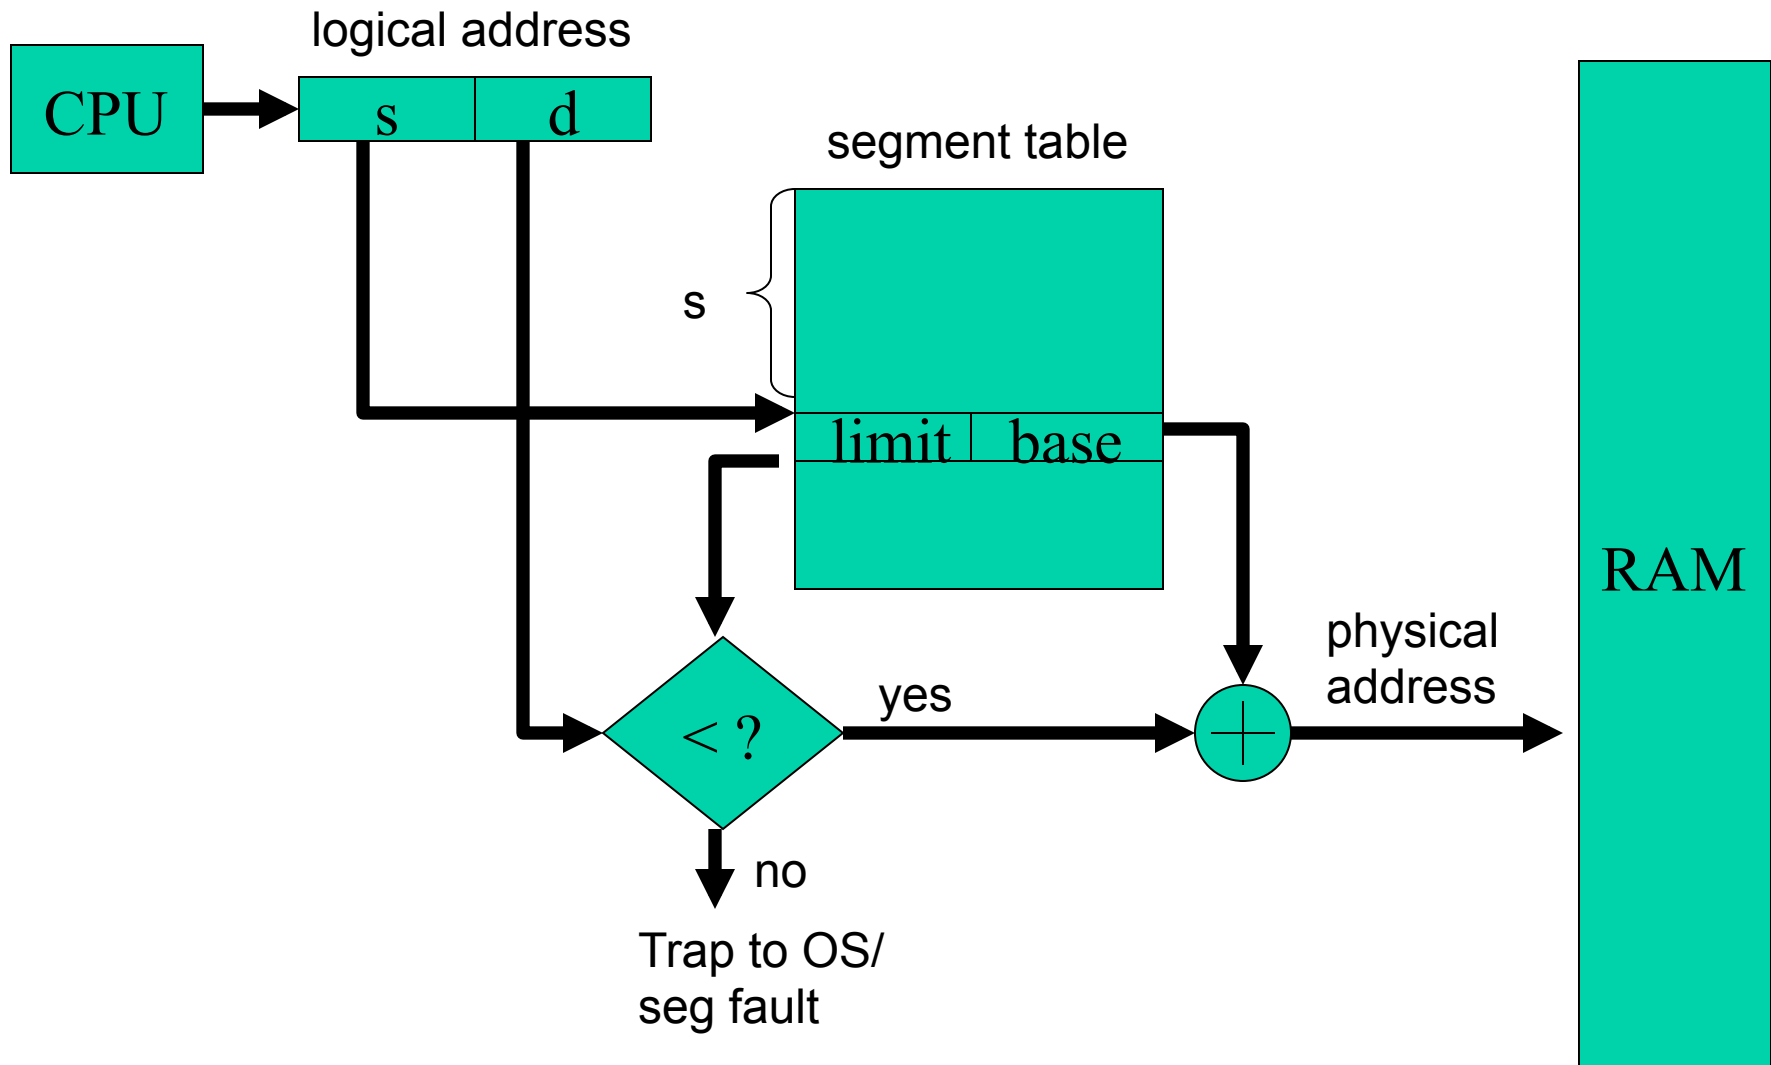
Segmentation

- An alternative to dealing with external fragmentation using fixed size pages is instead to employ variable sized *segments*
 - Subdivide a process into variably sized segments that are organized according to some logical criteria
 - a process has code, data, stack, and heap - each of these could be a separate segment
 - a process could subdivide its code into functional segments, e.g. a Web server could subdivide via its functional components into a networking segment, a database interface segment, and a dynamic page composition segment, etc.

Segmentation

- Each instruction in a segment has a logical address = <segment #, offset>
- Need a segment table to keep track of where each segment is mapped in main memory
- Each entry of the segment table needs a base and limit field to place in the base and limit registers, because segments are variably sized

Segmentation



Segmentation

- Pentium supports pure segmentation and segmentation with paging
- Linux on the Pentium uses segmentation sparingly
 - my reading of the literature suggests that paging is more popular than segmentation
 - Linux supports 3-level paging
- Drawback of segmentation: fragmentation of free space in RAM becomes even more complex and hard to manage